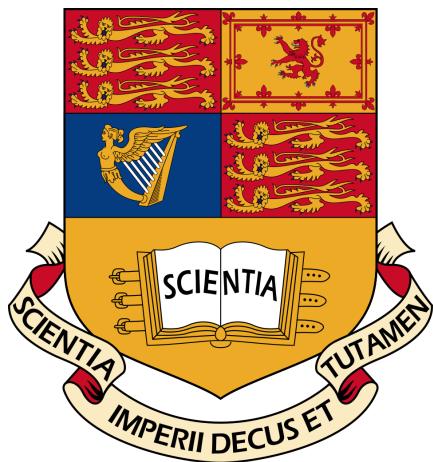


Imperial College London

Department of Electrical and Electronic Engineering

Final Year Project Report 2021



Project Title: **State-decomposition Techniques for Efficient Deep Reinforcement Learning**
Student: **Edoardo David Santi**
CID: **01373388**
Course: **4T**
Project Supervisor: **Professor Kin K. Leung**
Second Marker: **Dr Tania Stathaki**

Final Report Plagiarism Statement

I affirm that I have submitted, or will submit, an electronic copy of my final year project report to the provided EEE link.

I affirm that I have provided explicit references for all the material in my Final Report that is not authored by me, but is represented as my own work.

Acknowledgments

I would like to thank my supervisor Prof. Kin K. Leung for the continuous support and availability throughout the duration of this project and my excitement about the prospect of working with him as my PhD supervisor in the future. I would also like to express my gratitude towards my parents for their continuous support and encouragement throughout my studies.

Abstract

The existence of prohibitively large state spaces is a common problem of reinforcement learning algorithms in practical domains. This project proposes state-decomposition approaches to alleviate this problem by allowing more efficient learning in deep reinforcement learning algorithms. This exploits the almost independence of sub-systems within a large system that only rarely interact with each other. The developed state-decomposition methods rely on learning sub-agents for each of these sub-systems. This report focuses on grouping the states into subspaces which have low probability of transitions between each other. The developed algorithms are tested in "toy" environments and their performance was compared with a baseline Deep Q-learning algorithm. Amongst the developed techniques, two demonstrated better performance than the baseline algorithm. The algorithm named 'DeltaSwitch' achieved performance gains in both perfectly decomposable environments (perfectly separable into sub-systems) and more generic environments. The performance gain disappears once the sub-systems are not independent enough. Another algorithm that achieved performance gains is based on "progressive networks". These techniques were thus demonstrated capable of improving the efficiency of learning and a possible application of these techniques includes the allocation of processing jobs in inter-connected data-centres. Further studies can focus on the use of different state-decompositions, not based on transition probabilities and on techniques to share common information between similar subspaces.

Contents

1	Introduction	2
1.1	Motivation	2
1.2	Project Definition	2
1.3	Structure of the report	3
2	Background	4
2.1	Reinforcement learning	4
2.2	Artificial Neural Networks	8
2.3	Q-learning and Deep Q-learning	9
2.4	Problem of large state spaces	10
3	state-decomposition method	12
3.1	Originally proposed method	12
3.2	Decomposing the state-trans. matrix	15
4	Experimental method	17
4.1	Tools	17
4.2	Testing environments	18
4.3	Benchmark DQN agent	22
4.4	Ensuring fair comparison	23
5	Evaluating results	25
5.1	Evaluation method	25
5.2	Effect of state representations	26
5.3	Effect of NN size	28
5.4	Initial results	29
5.5	'DeltaSwitch'	30
5.5.1	Applying 'DeltaSwitch' to non-prefectly decomposable environment	35
5.5.2	Effect of using reduced encoding	37
5.5.3	Effect of wall	38
5.6	A different approach based on transfer learning	39
5.7	Techniques to learn the 'switch'	45
6	Applicability of the decomposition method	48
7	Conclusions and further work	52
7.1	Key findings	52
7.2	Discussion	53
7.3	Further work	54
A	Performance graphs	60

Chapter 1

Introduction

1.1 Motivation

Reinforcement learning (RL) has been a very active research area in the last few years, achieving impressive feats in a wide range of applications. Some advantages of reinforcement learning algorithms are great flexibility, the capability of solving both very low and high-level problems [28] and not requiring a model of the environment in which they are used. RL has famously been applied to games, such DeepMind's AlphaGo [26], a program that in 2016 beat the top-ranked Go player in the world. Other successful applications include healthcare [37], robotics [2], autonomous driving [7] and natural language processing [22]. My project supervisor Prof. K. K. Leung has also been working on applications of RL in software-defined networks (SDNs), such as for the control of the placement of services [38] and the synchronisation of controllers in distributed SDNs [39].

Dealing with very large systems is one of the main problems encountered in the application of reinforcement learning in practical environments. RL is based on learning through interaction with the environment. As the size of the system increases, so does the complexity of the problem and it becomes increasingly difficult to learn how to interact with the environment, often making the learning process excessively long. In practical scenarios, it might not be feasible to train an RL agent for long enough for it to be a feasible solution, as this might incur excessive costs or, in the case that the behaviour of the environment changes over time, the algorithm might not be able to learn fast enough to adapt to the changing environment.

1.2 Project Definition

The goal of this project is to develop a state-decomposition method, that aims to alleviate the problem of large state spaces of the Markov Decision Processes (MDPs) that underlay many practical environments. This new method leverages the characteristic of many control problems of being composed of many almost independent sub-problems that only seldom interact with each other. The method considers these sub-problems independently as well as combining them to form a global agent that controls the whole system and takes into account their interactions.

This project deals with developing reinforcement learning algorithms that use the high-level idea of state-decomposition and to test and benchmark them against baseline algorithms to determine

if such methodology can provide any performance gains. The performance is measured in terms of training samples or episodes, which is the amount of training experience required to achieve a certain level of accuracy by the algorithm. Although training time is usually related to the number of training samples used, this work does not directly consider the reduction of training time (or computational effort).

1.3 Structure of the report

This project's report has the following structure. Chapter 2 provides the background information required to understand the topic of this project. This is about reinforcement learning basics, neural networks, Q-learning and Deep Q-learning, as well as introducing the problem of large state spaces which this project tries to alleviate and methodologies used in the literature to address the issue. Chapter 3 introduces the state-decomposition method more specifically, explaining the high-level ideas and the originally devised state-decomposition method. Chapter 4 presents the tools and methods that are used in testing the state-decomposition technique and the type of experiments that are run. Chapter 5 shows the most important results obtained during testing and their interpretation. Chapter 6 discusses the applicability of the proposed techniques and provides a real-life example where the technique can be useful. Lastly, Chapter 7 discusses the findings of this project, its limitations and possible extensions.

Chapter 2

Background

2.1 Reinforcement learning

"The reinforcement learning problem is meant to be a straightforward framing of the problem of learning from interaction to achieve a goal" [28].

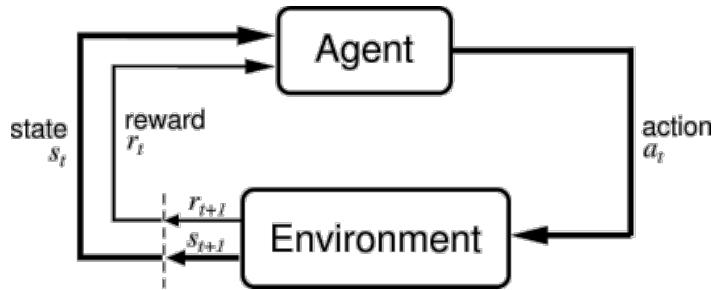


Figure 2.1: Interaction of the agent with the environment during time step t . Figure taken from [28].

The reinforcement learning *agent* continuously interacts with the *environment*. This interaction can be modelled as a sequence of time steps $t = 0, 1, 2, 3, \dots$ ¹. At each step t , the environment is represented by its *state* $S_t \in \mathcal{S}$, where \mathcal{S} is the set of all the possible states. Taking the *state* as input, the agent selects an *action* $a_t \in \mathcal{A}$, where \mathcal{A} is the set of all the possible actions. In the next time step $t + 1$, depending on the current state S_t and the chosen action a_t , the environment transitions to a new *state*² S_{t+1} and emits a *reward* $R_{t+1} \in \mathcal{R} \subset \mathbb{R}$, which the agent uses as feedback information to "learn" how to interact with the environment. Each time step can then be represented by the $\langle S_t, A_t, R_{t+1}, S_{t+1} \rangle$ tuple.

The agent implements a mapping from the possible states to the probabilities of choosing the different actions. This mapping is called *policy* and at time step t , it adopts symbol π_t . The probability of choosing action $A_t = a$ given the current state $S_t = s$ is denoted by $\pi_t(a|s)$. The goal of reinforcement learning is to learn the optimal policy, which is the policy that maximises the rewards, through

¹The interaction with the environment can also be modelled in continuous time, however this is outside the scope of this project.

²Note that it is possible that $S_{t+1} = S_t$

repeated interaction with the environment. In general, we aim to maximise a function of the rewards at all time steps rather than the immediate reward, so the optimal policy doesn't necessarily maximise immediate rewards. The simplest case is to maximise

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_T \quad (2.1)$$

where T is a final time step. This approach can be used when the agent-environment interaction has an end, in which case a sequence of interactions from the initial time to time T can be called an *episode*. In other cases, the interactions can continue indefinitely, thus this reward function is not feasible as $T = \infty$ and G_t could be infinite. To fix this, we apply a *discount rate* γ , with $|\gamma| < 1$, to future rewards:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}. \quad (2.2)$$

This allows G_t to have a finite value in indefinitely continuing tasks, which are those tasks that do not have finite time episodes.

In general, the dynamics of the environment could depend on the whole history of states, rewards and actions over an episode. However, we aim to have an agent which can choose the optimal action at each time step based only on the current state. Thus the knowledge of the current state and action should give the maximum possible amount of information about the dynamics of the environment. In this case, knowledge of the history of states, rewards and actions in the previous time steps does not provide any additional information about the dynamics of the environment than knowing the current state. States that follow this property are said to follow the *Markov property*. More formally, if the states follow this property

$$Pr\{R_{t+1} = r, S_{t+1} = s' | S_0, A_0, R_1, \dots, S_{t-1}, A_{t-1}, R_t, S_t, A_t\} \quad (2.3)$$

$$= Pr\{R_{t+1} = r, S_{t+1} = s' | S_t, A_t\} \quad (2.4)$$

$$= p(s', r | s, a). \quad (2.5)$$

This is the equation that governs Markov Decision Processes (MDPs) and hence a reinforcement learning task formulated with states that follow the Markov property is an MDP.

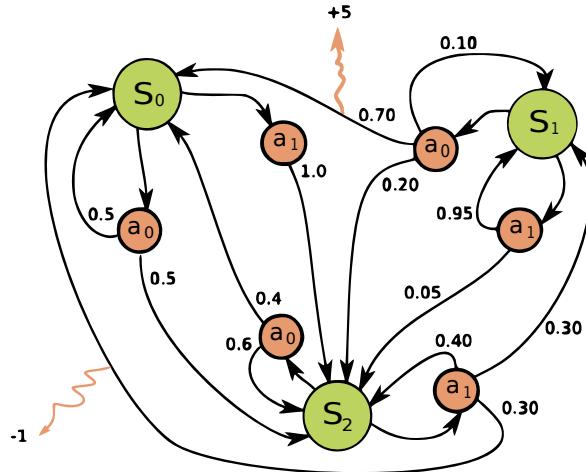


Figure 2.2: Graph representation of a Markov Decision Process. Note the rewards (zigzag arrows) that are released in certain transitions. [3]

Most reinforcement learning algorithms involve the estimation of the *value functions*, which are functions of states or state-action pairs. A value function represents how good it is to be in a certain state or to perform a certain action in a certain state, in terms of the expected return G_t and given that we are following a certain policy π . The action value function is defined for each pair of (s, a) with $s \in \mathcal{S}$ and $a \in \mathcal{A}$ as

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a] = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s, A_t = a \right]. \quad (2.6)$$

Similarly, the state value function is given by

$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s] \quad (2.7)$$

$$= \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right] \quad (2.8)$$

$$= \sum_a \pi(a|s) \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s, A_t = a \right] \quad (2.9)$$

$$= \sum_a \pi(a|s) q_\pi(s, a). \quad (2.10)$$

It is useful to specify the value functions in recursive format, known as the Bellman equations. This gives:

$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s] \quad (2.11)$$

$$= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} | S_t = s] \quad (2.12)$$

$$= \sum_a \pi(a|s) \sum_{s',r} p(s', r | s, a) [r + v_\pi(s')] \quad (2.13)$$

and

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a] \quad (2.14)$$

$$= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1}|S_t = s, A_t = a] \quad (2.15)$$

$$= \sum_{s',r} p(s',r|s,a)[r + v_\pi(s')] . \quad (2.16)$$

For finite MDPs it can be proved that there always exists an *optimal policy* that maximises the state values of all states the action values of all state-action pairs, giving the optimal state values $v_*(s)$ and $q_*(s, a)$:

$$v_*(s) = \max_\pi v_\pi(s), \quad \forall s \in \mathcal{S} \quad (2.17)$$

$$q_*(s, a) = \max_\pi q_\pi(s, a), \quad \forall s \in \mathcal{S}, \quad \forall a \in \mathcal{A} . \quad (2.18)$$

Given full knowledge of the dynamics of the system³ and a fixed policy, we are able to evaluate the optimal state values and action values using methods based on dynamic programming [9]. One such method called *iterative policy evaluation*, which adopts the update step based on the Bellman equations given below, can be shown to converge to the state values:

$$v_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left[R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s, s', a) v_k(s') \right] . \quad (2.19)$$

Policy iteration is one possible method of evaluating the state values when the optimal policy is not given, which alternates steps of policy evaluation and updates of the policy. The policy is updated so that in every state the greedy action, meaning the one with the highest action value, is always picked. This process of solving the MDP when given the model of the environment is known as *planning*.

However, in many applications the model of the environment is unknown. Given a certain policy, it is possible in this case to approximate the state and action values by sampling the interaction with the environment. Monte Carlo methods interact with the environment by following the given policy and approximate each state value using the returns obtained in each episode after the first or each visit to the state. The update that occurs every time the state occurs for the first time in the episode (or every time, depending on the version of the algorithm) is

$$V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)] \quad (2.20)$$

The disadvantage is that state values cannot be updated until the end of an episode. Temporal Differences methods are another class of methods that solve this problem. Instead of relying on returns (G_t), they approximate them as the sum of the reward and the discounted state value of the next visited state. The simplest TD algorithm's update step is given by

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t)) \quad (2.21)$$

There are many applications of reinforcement learning in which state and action spaces are continuous. In that case, the MDP is not finite as there is an infinite number of possible states and actions. This will not be discussed as it is outside of the scope of this project, which only deals with discrete state and action spaces.

³The probabilities of state-transitions and rewards given the current state-action pair.

2.2 Artificial Neural Networks

Artificial neural networks are a type of function approximators that have been very popular in the last years thanks to improvements in computational power which made them a favourite choice in many data-driven applications, such as pattern recognition, computer vision, etc... in which they achieve state-of-the-art performance. They are especially useful as they are able to 'learn' the underlying structure of very large dimensional data when fed enough training samples.

The building block of artificial neural networks is the *perceptron*. The *perceptron* learns the values w_i with $i = 0, 1, \dots, n$ so that it successfully learns how to classify the inputs into two classes. The *perceptron* applies the function $w_0 + \sum_i = 1^n x_i w_i$, where x_i for $i = 1, 2, \dots, n$ are the elements of the input vector. The output of this function is then passed through a step function which maps all values to two possible values.

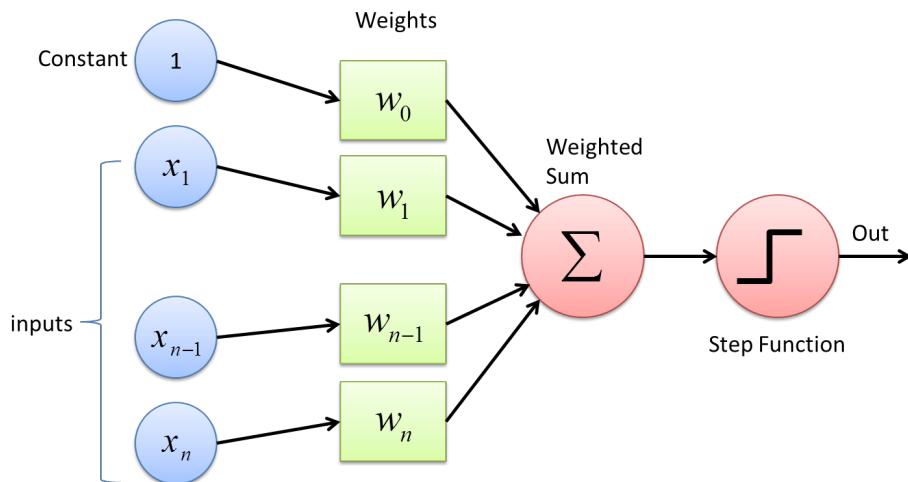


Figure 2.3: Graphical representation of a perceptron. Figure taken from [25].

Fully-connected neural networks are formed by combining multiple perceptrons together. By combining multiple perceptrons we are able to learn a function that maps from the inputs to multiple outputs. In general, fully-connected neural networks are formed by many layers, where each layer is formed by many perceptrons in parallel and the outputs of one layer are fed as inputs of the next layer. Fully-connected neural networks are called this way because all the outputs of a layer are fed as inputs to all the perceptrons of the next layer. By combining multiple layers together we are able to learn a wider range of functions than if we were only using one layer, making the neural network a more powerful and general approximator. While in the original perceptron we apply a step function to the intermediate function, in neural networks a variety of non-linear functions is used instead. This is called the activation function and it is useful as it allows the network to be more powerful by making it able to learn non-linearities. If there were no activation function, the neural network would simply learn a linear mapping between input and output, thus not being more powerful than a simple perceptron without step function.

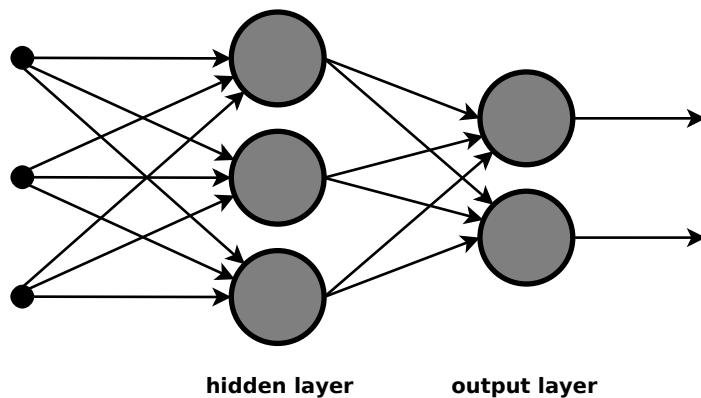


Figure 2.4: Fully-connected neural network.

As neural networks aim to learn a mapping function, they need to be trained using a training dataset. This is done by choosing a loss function, which penalises differences between what the network's outputs currently are for a given input sample in the dataset and its corresponding output in the training dataset. A common example of loss function is the mean square error, given by $\|\hat{y} - y\|^2$, where \hat{y} is the current output of the neural network while y is the "correct" output in the dataset. The most basic way to train a neural network is to iterate through the training samples multiple times, and for each of them adjust the weights of the perceptrons in the network in order to reduce the loss function. This can be done by using gradient descent, however nowadays there are many different optimisers that have shown to have better rate of convergence. Running optimisation steps using single sample leads to high variance, thus we normally run optimisation steps based on the loss function of batches, that is each optimisation step aims to minimise the loss function for a batch of samples rather than single samples.

2.3 Q-learning and Deep Q-learning

Many RL algorithms deal with the problem of solving an MDP without being given the optimal policy and the dynamics of the system. Q-learning [35] is one of the most noteworthy among these. It is based on Temporal Differences learning as the action value updates follow a very similar formula. The basic value update in Q-learning is given by:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, a)] \quad (2.22)$$

Differently from the previously described TD algorithm, this operates on action values instead of state values. In Q-learning, the value update is done using the term $\max_a Q(S_{t+1}, a)$, which is the action value of the greedy action in the next visited state. This is not necessarily the action taken in the following step when visiting S_{t+1} , thus this an *off-policy* algorithm.

Differently from the Temporal Differences algorithms that aim at evaluating a fixed policy, Q-learning also deals with the control problem of defining the policy. Q-learning is normally programmed so that its current policy is based on its current action value estimates. An example is the ϵ -greedy policy, in which at each time step, the greedy action is picked with probability $1 - \epsilon$ and a random action is picked with probability ϵ .

This algorithm is only defined for discrete environments and it is proven to converge to the optimal action values with a probability of 1, as long as all the states and actions are visited multiple times [34]. This requirement may not be feasible when the number of states is very large and it is impossible when the number of states is infinite such as when the MDP's state space is continuous. Using a learned Q-function approximator instead of simply storing the action values in tabular form can solve this issue, as it is possible to learn a mapping function that generalises over the whole state space without having to visit all the possible states.

Deep Q-Learning (DQN) [17] implements the Q-function approximator using a deep neural network. This was demonstrated to achieve state-of-the-art performance in learning to play Atari games directly from the pixel data [16]. DQN takes advantage of the ability of deep learning to learn from high-dimensional data representations to be able to directly feed large state spaces into the agent. Deep neural networks rely on the assumption that they are fed independent samples, however in reinforcement learning the samples are correlated as they are generated by a temporal sequence. An *experience replay* is used to solve this. The transition samples are saved in memory and a random batch is selected to train the neural network at each time-step, thus removing correlations due to temporal proximity of the samples. This also allows the update to be based on a batch, rather than a single sample, which improves the sample efficiency⁴.

One issue of the Q-learning algorithm is that using the maximum action value of the next visited state in the update step results in a positive bias in the estimation of the action values, which can lead to poor performance in some stochastic environments. Double Q-learning [36] deals with this problem by splitting the samples randomly between two separate sets of action value estimates and by updating the values of each set of estimates using the values of the other, which can be proven to remove the positive bias⁵. The same problem persists in deep Q-learning and a similar solution has been devised in Double Deep Q-Learning [32]. Another version of DQN that achieves improved performance in typical benchmarks such as playing Atari games is Dueling DQN [33], which operates on the concept of expressing action values as a sum of state values and advantages: $\mathcal{Q}_\pi(s, a) = \mathcal{V}_\pi(s) + \mathcal{A}_\pi(s, a)$.

While the original Q-learning algorithm can only be applied to systems with discrete state and action spaces, other algorithms based on Q-learning have been proposed to deal with continuous state and action spaces, such as Wire Fitted Neural Network Q-learning [12].

2.4 Problem of large state spaces

Very large systems correspond to MDPs with a very large number of states, which can be difficult to solve using reinforcement learning algorithms. The number of possible state-transitions is the square of the number of states, thus the complexity of controlling a system grows very quickly as the size of a system increases, and learning an optimal policy becomes very difficult and often infeasible.

⁴How efficiently the samples are used by the algorithm to learn. A more sample efficient algorithm requires fewer interactions with the environment to learn an optimal policy.

⁵It actually introduces a small negative bias, thus tending to underestimate the action values.

Various approaches have been used to solve this issue. One general approach is to learn representations of the states to improve generalisation to unseen states or to decrease the dimensionality of the states by compression. DQN [16] exploits the ability of deep learning to learn from high-dimensional data and to generalise to solve this issue, though this is often not sufficient. [23] jointly learns state and action embeddings to improve generalisation. In World Models [13], a compressed latent representation of the states is explicitly learned by an autoencoder, which is then fed into the reinforcement learning model instead of the original states. Other methods rely on state aggregation [10], in which states with similar state transition probabilities and rewards are grouped to form a higher-level state. This technique can also be applied to hierarchical reinforcement learning to define the states of the global learner [6].

Hierarchical reinforcement learning [8] is another approach that reduces the impact of large state spaces. Hierarchical RL provides temporal abstraction by subdividing the main task into sub-tasks. For example, the task of commuting to work can be subdivided into sub-tasks such as opening a door, turning right, walking to the tube station, etc. The agent is composed of hierarchies of sub-agents in which each sub-agent is fed a compressed or reduced state space in input and outputs high-level actions to its "children" sub-agents that perform the lower-level actions required to complete the received higher-level action. Examples of hierarchical reinforcement learning algorithms are based on MAXQ value decomposition [11], Options [29] and Hierarchies of Machines (HAM) [21].

Model-based reinforcement learning [18] has been receiving increasing attention lately. This class of algorithms is at the intersection of planning algorithms and model-free reinforcement learning as it approximates the model of the environment during the learning process, which can be used to more efficiently learn from new samples. Dyna [27] learns a model of the environment and uses it to simulate samples for Q-learning other than those obtained by actual interaction with the environment. AlphaZero [26], which achieved state-of-the-art performance in playing Chess, leverages full knowledge of the model⁶ to perform a Monte-Carlo tree search for every step in the real-world environment. [13] trains a recurrent neural network model that predicts the following state at each-time step. Though this class of algorithms does not explicitly tackle the problem of large state spaces, the improvements in sample efficiency can still solve the issue by allowing to train a model with fewer interactions with the environment.

⁶The model is given at the start rather than learnt.

Chapter 3

state-decomposition method

The main topic of this project is the development of the state-decomposition method for reinforcement learning. This is suited to control problems of environments that are composed of many sub-problems that only seldom interact with each other. The state-decomposition method takes advantage of the almost independence of these problems to learn agents that deal with these smaller problems, which are easier to solve. These separate agents need to be combined to take into account the interactions between the sub-problems.

An example of a system composed of many almost independent sub-systems could be the control of a communication network composed of many networks that rarely communicate with each other. The state-decomposition method would aim to learn an optimal controller by training agents that control the smaller sub-networks and combining them to account for the small amount of interaction between sub-networks.

3.1 The original state-decomposition method

This section discusses the procedure of the initially proposed state-decomposition method. This is not the only version that was considered in this project, as the state-decomposition can be regarded as a more general idea that can be exploited in different ways.

The first step of the process is to separate the sub-problems by identifying separate groups of states. This is done by examining the state transition probability matrix. Given a finite MDP, this is defined as the matrix where each entry P_{ij} is the probability of transitioning from state $S_t = i$ to the next state $S_{t+1} = j$:

$$\begin{bmatrix} P(S_{t+1} = s_0 | S_t = s_0) & P(S_{t+1} = s_1 | S_t = s_0) & \dots & P(S_{t+1} = s_n | S_t = s_0) \\ P(S_{t+1} = s_0 | S_t = s_1) & P(S_{t+1} = s_1 | S_t = s_1) & \dots & P(S_{t+1} = s_n | S_t = s_1) \\ \dots & \dots & \dots & \dots \\ P(S_{t+1} = s_0 | S_t = s_n) & P(S_{t+1} = s_1 | S_t = s_n) & \dots & P(S_{t+1} = s_n | S_t = s_n) \end{bmatrix} \quad (3.1)$$

where n is the number of states in \mathcal{S} .

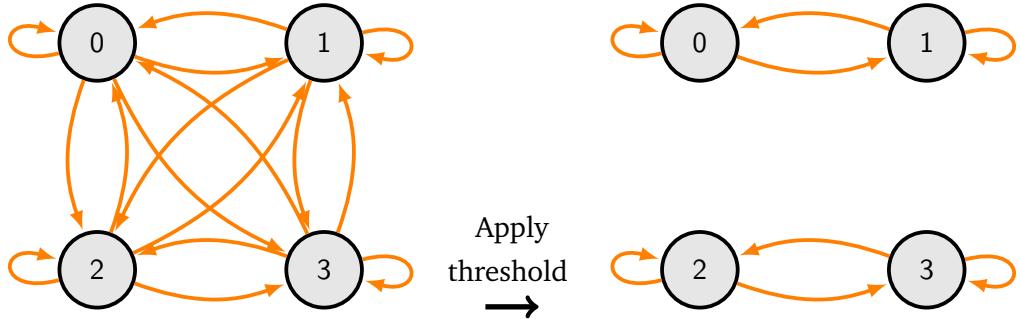


Figure 3.1: Applying a threshold to the state transition matrix and normalising the probabilities can cause the MDP to separate into separate MDPs as it reduces the number of possible transitions.

The state-decomposition method is based on the fact that some state-transition probabilities are negligible compared to others and could be regarded as 0 to obtain a sparse version of the state-transition matrix. Specifically, every matrix element smaller than a certain threshold τ can be set to 0. The probabilities in the matrix can then be normalised so that each row of the matrix has a sum of 1. This step is unnecessary and can be skipped in the implementation; however, it is necessary to obtain a valid MDP from the state-transition matrix after applying the threshold. Following this procedure, the original MDP could have transformed into multiple independent MDPs with smaller state spaces, as shown in Figure 3.1. In this case, by reordering the states it is possible to express the state-transition matrix as a block matrix in the form:

$$\begin{bmatrix} B_1 & 0 & \dots & 0 \\ 0 & B_2 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & B_n \end{bmatrix} \quad (3.2)$$

where the matrices B_1, B_2, \dots, B_n are the state-transition matrices of the newly formed separate MDPs. This is the method of state-decomposition that this report focuses on, which is based on ignoring state transitions with probability lower than the threshold τ . There are different ways to decompose the state space, the choice of which might depend on the specific problem that the method is applied to. This is discussed further in Section 7.3.

The following describes the original architecture that is devised to apply state-decomposition to a practical problem. Once the states of the original MDP have been separated into multiple subspaces, it is possible to proceed with training the RL agent. In the first stage of training, a separate DQN agent is trained for each of these subspaces. To do so, we consider for each subspace's agent only those state transitions that start and end within the corresponding subspace. Once the subspaces' agents converge, these should be combined to form a global agent that also takes into account transitions between the different subspaces. This is done by feeding the action values of the subspaces' agents, which are the outputs of their neural networks, as inputs of another neural network, as in Figure 3.2. This forms a bigger neural network which is the Q-function approximator of a new global DQN agent, which is then trained using all possible transitions. This will be referred to as the second stage of training and the method as a two-stage method.

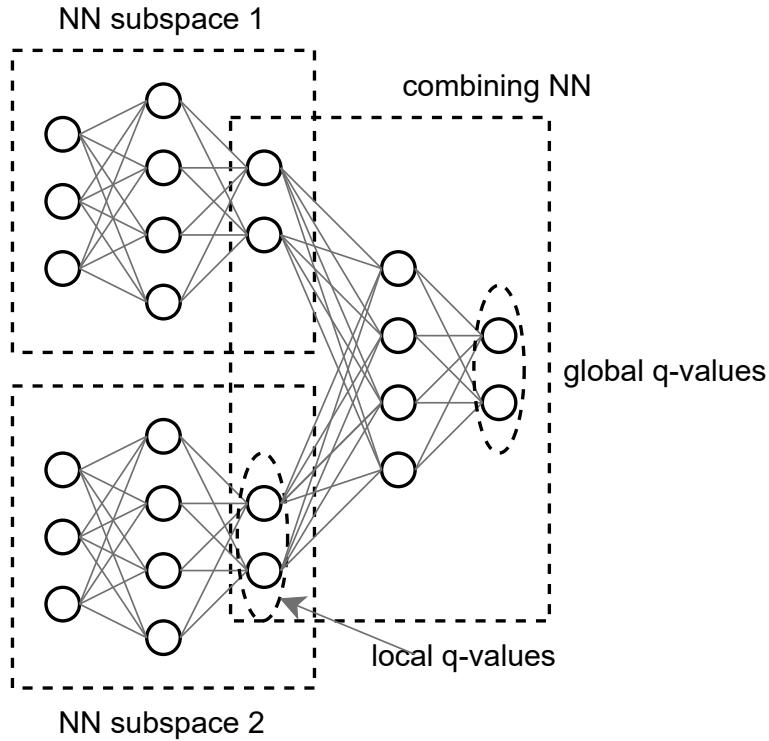


Figure 3.2: The NNs of the decomposed states' agents are merged as input of another NN to form a new bigger NN architecture.

This method operates on the assumption that it would require fewer samples to achieve convergence for RL agents operating on many smaller MDPs than for one single MDP having as the number of states the sum of the number of states of the smaller MDPs. By then combining the neural networks as previously explained, we take into account interactions between the sub-problems, allowing us to achieve the optimal policy for the original MDP. For this method to be successful, it is assumed that the weights of the neural networks of the subspaces' agents after the first stage are not far from the values they take once the global agent converges in the second stage. This is equivalent to saying that the first stage of training provides a very good initialisation of the weights of the network used in the second stage. Although we use DQN agents, this method could be generalised to any agents that use a function approximator to evaluate the action-values given the states as input. This method however would not be as useful if no approximator is being used and the action-values are simply stored in a table, such as in the original Q-learning algorithm. In this case, applying the state-decomposition would simply ignore the transitions between different subspaces without changing the structure of the agent at all in the first stage of training, leading to no benefits at all.

It is important to mention that the state-transition matrix depends on the policy being followed by the agent that interacts with the environment. Thus the state-transition matrix is most likely different when starting to train an agent and upon convergence. For this reason, different state-decompositions could be possible depending on the policy that the state-transition matrix is based on. For example, some state-decompositions might exist regardless of the policy, such as when the transitions between two subspaces are made unlikely by the characteristics of the environment itself, such as moving between two places separated by a wall in a navigation problem. However, some

decompositions are purely created by the policy of the agent, such as if in a navigation problem there is no wall separating two positions, but the agent simply never moves between these two places. More examples are given in the next chapter and for now it is sufficient to acknowledge that different decompositions can occur depending on the assumed policy.

This method is the one that was initially proposed to study the effects of state-decomposition. In chapter 5 we show that this architecture is problematic, having issues such as the first stage of training not successfully providing a good initialisation for the second stage and thus not speeding up the training of the second stage.

3.2 Decomposing the state-transition matrix

This report focuses on the application of state-decomposition where the state-decomposition is "intuitive" and it is done in a non-mathematical way, as will be shown in the next chapter. However this is not always the case, thus we need a systematic way to obtain the correct state-decomposition, given that we know the state-transition matrix. This would normally also be unknown and techniques to estimate this could be investigated as an extension of this project.

The decomposition into subspaces of states given the state-transition table can be done using Algorithm-1, which takes in input the state transition matrix. The states are initialised as all separate subspaces. Then we iterate through all transition probabilities and given that the transition probability from state i to state j is higher than τ we join the subspaces containing the two states. After iterating through all transitions this returns the decomposed subspaces.

Algorithm 1: State-decomposition algorithm.

```

stateTransMatrix: input 2D array representing the state-transition matrix
subspaces : output data structure representing the decomposed subspaces

subspaces = initialise as one subspace for each group;
for  $i \leftarrow 0$  to  $height(stateTransMatrix)-1$  do
    for  $j \leftarrow 0$  to  $width(stateTransMatrix)-1$  do
        if  $stateTransMatrix[i][j] > \tau$  then joinSubspacesContaining( $i, j$ );
    end
end
```

Sometimes we know the number of subspaces that the original MDP should be decomposed into. For example, we might know the characteristics of the system and the number of subsystems it can be considered to have, such as how many sub-networks compose a certain communication network. In such case, the function that returns the decomposed subspaces can be run repeatedly while updating the threshold value τ at each iteration using binary search, until the MDP is decomposed in the correct

number of subspaces. This is given in Algorithm-2.

Algorithm 2: State-decomposition into given number of subspaces

```

stateTransMatrix : input 2D array representing the state-transition matrix
nSubspacesDesired: input integer representing the number of subspaces to decompose into
subspaces : output data structure representing the decomposed subspaces
maxIterations  $\leftarrow$  50;
changeFactor  $\leftarrow$  10;
 $\tau \leftarrow 1.0$ ;
upperBound  $\leftarrow$  NONE;
lowerBound  $\leftarrow$  NONE;
for  $i \leftarrow 1$  to maxIterations do
    if upperBound and lowerBound are defined then
        |  $\tau \leftarrow (\text{upperBound} + \text{lowerBound})/2$ ;
    end
    if only upperBound is defined then  $\tau \leftarrow \text{upperBound}/\text{changeFactor}$ ;
    if only lowerBound is defined then  $\tau \leftarrow \text{lowerBound} * \text{changeFactor}$ ;
    subspaces  $\leftarrow$  stateDecomposition (stateTransMatrix,  $\tau$ );
    if nSubspaces(subspaces) == nSubspacesDesired - 1 then break;
    if nSubspaces(subspaces) < nSubspacesDesired - 1 then lowerBound =  $\tau$ ;
    if nSubspaces(subspaces) > nSubspacesDesired - 1 then upperBound =  $\tau$ ;
end
```

Chapter 4

Experimental method

This section describes the experimental method used to assess the performance of the state-decomposition method. This is done by implementing different versions of the state-decomposition method, testing them in different environments and comparing the performance with a baseline algorithm.

4.1 Tools

Python The programming language chosen for this project is Python. This is motivated by Python being a high-level language that allows writing compact, readable code and by the widespread availability of resources and packages for it. Additionally, Python can be run on Colab notebooks, which are described in the next paragraph.

Google Colaboratory (Colab) Colab is a platform from Google which is widely used for machine learning research and prototyping of new models. Colab combines a notebook platform that many Python developers are familiar with and the high-performance of cloud computing. The user can choose to run a program using CPUs, GPUs, or TPUs¹. Examples of the available GPUs are Nvidia K80 and NVidia P100D. The disadvantage of Colab compared to cloud computing platforms such as Google Cloud and Amazon AWS is that the user is not guaranteed full access to the hardware accelerators (such as GPUs) and thus the performance can vary. The session length is also limited to 12 hours. These problems are mitigated by purchasing a "Pro" license for the price of \$10 per month, which prioritises access to GPUs and increases the maximum session length to 24 hours. The main advantage is an easy to use platform thanks to its notebook format which allows for easy prototyping, rather than the less user-friendly access via SSH in the Terminal required by other cloud-computing platforms.

Keras In this work the deep-learning models are implemented using the deep-learning API Keras [31]. This is a high-level interface that uses Tensorflow [5] for its backend. Keras allows to code deep-learning models in a very readable and compact way. Having Tensorflow as its backend, it can run on CPUs, GPUs, and TPUs, taking advantage of the high-performance hardware offered by Google Colab.

¹Tensor Processing Unit: a custom hardware accelerator developed by Google to accelerate the linear algebra operations that occur in machine learning. [1]

OpenAI Gym As this project deals with the development and evaluation of a new method for reinforcement learning algorithms, it is decided to apply it in a "toy environment" that would simplify experimentation and provide insight into how the decomposition techniques would operate. OpenAI offers an open-source Python toolkit called "OpenAI Gym" [19], which is currently the most commonly used toolkit to benchmark reinforcement learning algorithms. This toolkit offers a variety of different environments, ranging from classical control problems such as "CartPole-v0" [4], to more complicated environments such as playing Atari games using the pixel data as sensory input or solving robotic manipulation problems. The "OpenAI Gym" offers environments with a wide range of complexities and all combinations of discrete and continuous state and action spaces. Figure 4.1 shows a screenshot of one of the most popular environments in the toolkit.

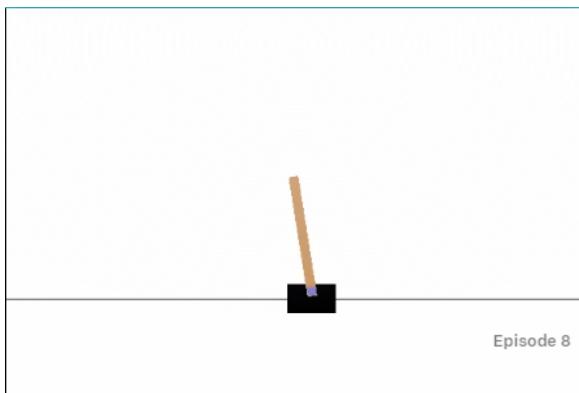


Figure 4.1: The CartPole-v1 environment is one of the most well-known environments of the OpenAI Gym library in which the goal is to balance the cartpole while keeping it within the boundaries of the screen by applying a leftwards or a rightwards force at each time-step. Screenshot of [4].

4.2 Testing environments

The state-transition matrices and their decomposed versions were obtained² for 4 different OpenAI Gym environments: "Cartpole-v1"³, "FrozenLake8x8-v0", "CliffWalking-v0" and "Taxi-v3". The state-decomposition algorithm is aimed at environments whose state spaces decompose into similarly sized subspaces, as this is the case in which the method would be potentially very effective. If one of the state spaces is much greater than the others and similar in size to the original state space, the reduction in complexity of training an agent for this subspace would not be as significant. Having subspaces that are too small (in the extreme case having just one state) limits the amount of learning that can be achieved by the subspaces' agents, so that learning to combine the sub-agents could be very complex and slow. The MDPs of "Cartpole-v1", "FrozenLake8x8-v0" and "CliffWalking-v0" don't decompose into similarly sized subspaces, thus they are not suitable for testing the state-decomposition method.

The "Taxi-v3" environment, shown on the left in Figure 4.2 has then been deemed to be a more suitable candidate. The environment has:

²These were obtained by "cheating" by simply inspecting the source code of the environments.

³This environment has continuous states which are discretised using a linear quantiser to apply the decomposition algorithm.

- 500 discrete spaces, corresponding to 25 taxi positions in the 5x5 grid, 4 possible destinations, and 5 possible passenger positions (1 of which is inside the taxi).
- 6 discrete actions: 4 moves (up, down, left, right), pick-up passenger, and drop-off passenger.
- Rewards: +20 for a successful drop-off, -1 at each time step (also when hitting a wall), and -10 for illegal pick-up and drop-off actions.

An episode terminates when the passenger is successfully dropped off or after 200 time-steps. Applying the state-decomposition to this environment can form 4 equally sized subspaces even with a threshold of 0, meaning that these are 4 completely independent subspaces that can be treated as independent Markov Decision Processes. These 4 subspaces correspond to the 4 possible destinations of the passenger, which never change during a single episode. Thus, this is a more specific and easier problem than the one we are trying to solve, where the subsystems, though rarely, do interact with each other.

This project then proposes a modified version called "TaxiTraps", shown in the middle in Figure 4.2, in which the transition probabilities were modified to allow a small amount of interaction between subspaces⁴ in order to test in the more general case. This has added "traps" located on the edges of certain cells of the grid. These activate with a low probability p_{trap} when the Taxi moves over them. If a trap does not activate, the taxi simply moves over it as in the original "Taxi-v3" environment. When the trap does activate, the behaviour of the environment is modified in two ways:

1. The destination of the Taxi changes, thus transitioning to a different subspace.
2. The agent receives a highly negative reward. This is added to ensure that the optimal policy in the "TaxiTraps" environment is different than in the original Taxi environment. Since in the first stage of the original state-decomposition method the transitions between different state subspaces are ignored, the agents are trained in an environment equivalent to the original "Taxi-v3". If the optimal policy is the same in "Taxi-v3" and "TaxiTraps", the optimal policy would not change between the first and second training stage of training, which is unlikely to happen in a general scenario. Thus, ensuring that the optimal policy changes between training stages is required for the results of testing to generalise to other environments.

During testing it has been noted that the two 'Taxi' environments would take a very high number of iterations to train and would often get stuck for long periods of time at reward -200, which is when the agent has learned to not commit illegal moves but it has not learned to drop-off the passenger successfully. Due to the limited computing power of my tools, this is a problem as it does not allow for significant testing with a high number of trials. Thus, it has been decided to try using a simpler version of the 'Taxi' environment, called 'Taxi2', shown on the right in Figure 4.2. The number of destinationations is reduced to 2 and the multi-stage nature of the task, having to pick-up the passenger, navigate the grid and then drop it off, is eliminated by setting the passenger to always be inside the Taxi, so that the pick-up and drop-off actions are eliminated and the task is completed by simply reaching the destination.

⁴Making transition probabilities between spaces of different subsystems non-zero.

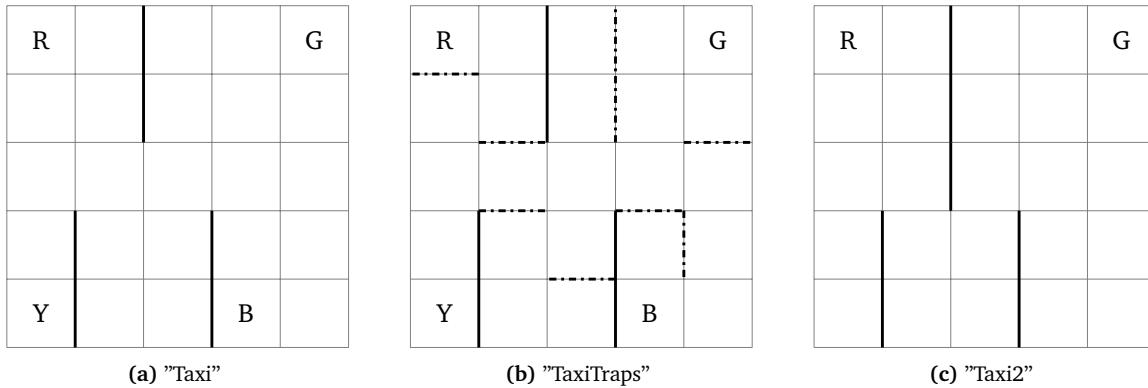


Figure 4.2: Original and modified 'Taxi' environments. The thick black lines represent the walls, while the dotted lines represent the "traps".

While 'Taxi2' reduces the complexity of training, this was still not simple enough and it didn't allow for easy interpretability of the results. I then introduced the 'Grid' environment, which is a simple 25 by 25 grid with two destinations, as shown in Figure 4.3. Those positions in the grid on the upper-left side of the dotted line in the figure are referred to as the 'upper-left triangle' or 'triangle 1', while those on the lower-right side are the 'lower-right triangle' or 'triangle 2'. At the start of the episode, the taxi is initialised in a random position, while the passenger is always assumed to be inside the Taxi, as in 'Taxi2', so that the pick-up and drop-off actions are eliminated and the learning process only deals with learning how to navigate in the grid. The two destinations are located at opposite corners and the initial destination is always chosen to be the one closest to the initial position of the taxi. 'Destination 1' is at the upper-left corner and 'destination 2' is at the bottom-right corner. Throughout an episode the destination always remains the same as in the initial state and the episode terminates when the taxi reaches the correct destination. To summarise, the environment has:

- 1250 discrete states, corresponding to 625 taxi positions in the 25x25 grid and 2 possible destinations.
- 4 discrete actions: move up, down, left and right.
- Rewards: +20 for being in the position chosen as the destination, -1 at each time step (also when hitting a boundary and thus not moving).

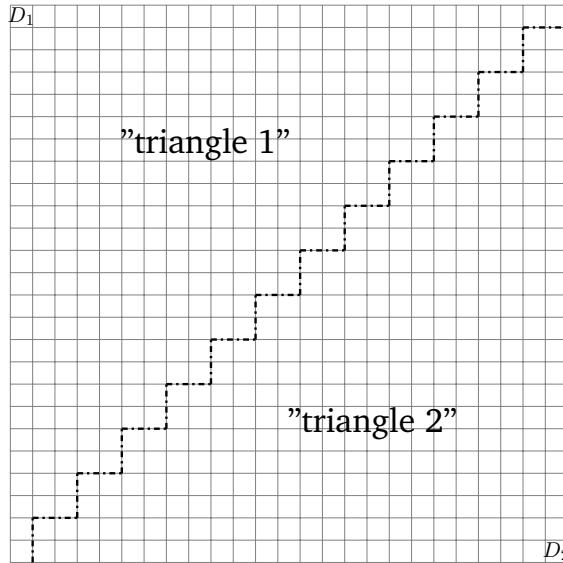


Figure 4.3: 'Grid' environment. The dotted lines have no effect on the environment, they simply separate the grid positions part of 'triangle 1' and 'triangle 2'.

This is a perfectly decomposable environment, as there are no possible transitions between states having different destinations. For a complete evaluation of the state-decomposition method, it needs to be tested in a more generic environment, with non-zero albeit small transition probabilities between the decomposed subspaces. This is achieved by introducing a probability ϵ of changing the destination whenever the current destination is the one closest to the current position in the grid. The initial destination can now be the far one, with probability ϵ . This modified 'Grid' environment will be referred to as 'GridEps'. Table-4.1 summarises the transition probabilities in 'GridEps'.

		next state	
		destination 1	destination 2
state	triangle 1 - destination 1	$1 - \epsilon$	ϵ
	triangle 1 - destination 2	0	1
	triangle 2 - destination 1	1	0
	triangle 2 - destination 2	ϵ	$1 - \epsilon$

Table 4.1: State transition probabilities in the 'GridEps' environment.

The transition probabilities are chosen in this way to achieve the following effects. That is, while providing small transition probabilities between different subspaces, changes in ϵ also change the optimal policy. This is important as the case in which the optimal policy is unaffected by the value of ϵ is a special case which should not be used for testing.

Another version that is used in various testings is very similar to 'GridEps' and will be called 'GridEps2'. The difference is in the probabilities with which the destination is picked in each time-step. In this case, when the current destination is the closer one between the two possible destinations, the destination in the next state is the closer destination to the position in the next state with probability $1 - \epsilon$ and the other one with probability ϵ , whereas in 'GridEps' the destination in the next

state depends on the position of the current state. This only makes a difference in transitions that cross the boundary between the two "triangles". In 'GridEps', when the boundary is crossed and the destination was the close one, the destination most likely remains the same, while in 'GridEps2' it most likely switches to the other one, which is the closest after crossing the boundary. This may seem like a small difference, but it has a significant impact on the effect of how the states are decomposed, as mentioned in the next chapter.

'GridEps' is an example of a scenario in which the decomposition is "intuitive", as mentioned in Section 3.2. It is clear here that one possible decomposition is given by creating one subspace for each possible destination, as at each-time the probability of changing destination is very small, being either 0 or ϵ . Another kind of decomposition that has been tested, although it does not necessarily follow the procedure previously described for decomposing the state-transition matrix for $\epsilon > 0$, is to assign states corresponding to positions in the upper-left "triangle" of the grid to one subspace and positions in the lower-right "triangle" to the other subspace. The reasoning is that when ϵ is 0, following the optimal policy would lead to never crossing the boundary between these two "triangles", meaning that they form two independent subspaces. When $\epsilon > 0$, this can happen and the transitions between the two subspaces are not small, thus it is an interesting example of using a different kind of decomposition that is not based on single transition probabilities.

It is possible to use policy iteration to determine the optimal policy and the optimal state values as we have access to the models' dynamics. This is used to make sure that the tested reinforcement learning algorithms, which are model-free, as they are not given the dynamics of the model as input, actually perform optimal actions and reach convergence after enough training.

4.3 Benchmark DQN agent

The aim of this project is to determine whether state-decomposition techniques can be used to increase the sample efficiency of reinforcement learning algorithms. Doing so requires comparing the performance of an algorithm where state-decomposition is applied and that of an equivalent algorithm without state-decomposition, while choosing all other parameters to be the same or to be the values that make for the fairest comparison. The number of possible choices to be made in these algorithms is such that providing a fair comparison can be a non-trivial task, as discussed in Section 4.4.

The chosen benchmark RL agent is a basic DQN agent. The value updates occur with discount factor $\gamma = 0.95$. It would be possible to use $\gamma = 1$, as the algorithm is tested in environments with finite length episodes, however, using a slightly smaller discount factor enhances the stability of the algorithm, giving more consistent results. The agent acts according to an ϵ -greedy policy. This involves a parameter called ϵ that is a different parameter and unrelated to the ϵ that represents the probability of transitions between different subspaces in 'GridEps'. This means that at each agent-environment interaction, the agent picks the action that currently has the highest estimated action value with probability $1 - \epsilon$ or a random action with probability ϵ . ϵ is initialised at 1, meaning that in the first iteration the action is completely random. ϵ decays by a factor 0.999 after each update of the action values. ϵ is set to stop decreasing at $\epsilon = 0.01$, this is to ensure that the algorithm does not stop exploring by having too small an ϵ .

Figure 4.4 shows the neural network that approximates the Q-values, which is fully connected and has 4 layers. The input state is fed as an integer to an encoding layer⁵, which then feeds to the first fully-connected layer, consisting of 512 neurons with "ReLU"⁶ activation function. This is then followed by 2 similar layers, having 256 neurons each. The last layer consists of a number of neurons corresponding to the number of actions in the environment where the agent is being used. Each output of the network corresponds to an action value of the input state, thus the activation function in the last layer is an identity function. This is to allow the outputs to take any possible values, without upper or lower bounds so that the network is theoretically capable of mapping any state to any possible action value⁷. The network is updated using "Minimum Square Error" as its loss function, that is, at each update step, the network weights are updated to reduce the function $L = \sum_{i=1}^n (y_i - f_\theta(x_i))^2$ where x_i are the inputs, y_i the corresponding outputs and f_θ is the function implemented by the neural network with weights θ . The network updates are done using the optimiser Adam [14], a commonly used optimiser that achieves better performance than gradient descent in many applications through the use of moments.

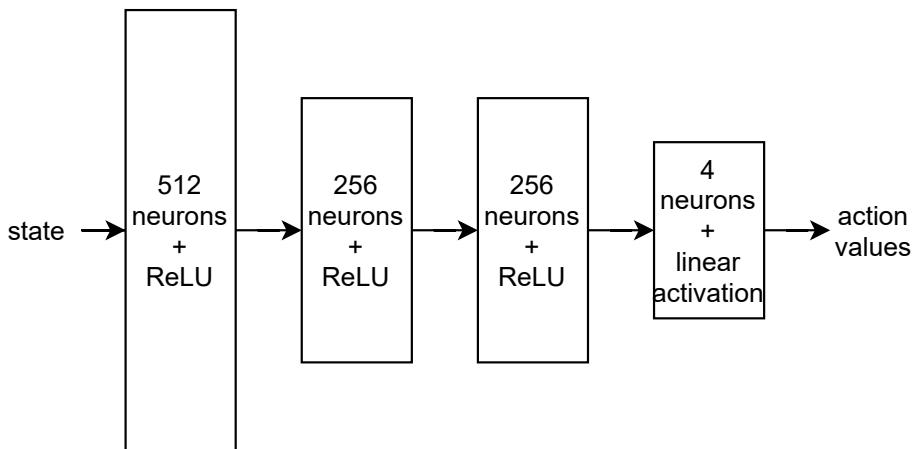


Figure 4.4: The neural network used for Q-values approximation in the baseline DQN algorithm.

Each iteration with the environment is saved in a replay memory as a $\langle S, A, R, S' \rangle$ tuple. The maximum size of this is set to be big enough so that all samples are saved, without having to discard them in the later stages of training. After each interaction with the environment, one update step of the neural network is run using a batch of 128 samples, randomly picked from the replay memory.

4.4 Ensuring fair comparison

When comparing the performance obtained using the state-decomposition technique with a benchmark method that does not use it, we would like ideally to isolate the effect of the state-decomposition effect, independent to other changes to other parameters. Unfortunately, it is not possible to perfectly isolate this effect, due to the complexity of the relationship between various hyperparameters. However, it is possible to make reasonable choices such that the comparison is as fair as possible. In this

⁵This is fixed and pre-programmed and does not include any learned parameters.

⁶"ReLU" stands for Rectified Linear Unit. It represents the function $f(x) = \max(0, x)$.

⁷Contrarily to for example a "ReLU" activation function, which wouldn't allow to have negative outputs

project, all hyperparameters are kept constant in all trials. Different experiments that involve different network architectures, such as using 2 neural networks instead of 1 always ensure that the total number of trainable parameters in the neural networks of each global agent is approximately equal. Other factors such as the encoding method of the state are always kept equal.

Chapter 5

Evaluating results of the proposed method

This chapter presents the most important results obtained during this project, using different variations of the state-decomposition method and comparing them with the baseline DQN algorithm.

5.1 Evaluation method

This project mainly deals with evaluating the sample efficiency of the state-decomposition method, which indicates how many experience samples are needed to attain a certain level of performance. During testing, the various agents were trained in the environments for multiple episodes. The sample efficiency affects how quickly the algorithm can learn an optimal policy and it can be illustrated by plotting the total reward obtained in the corresponding episode versus the number of experienced samples, as in Figure 5.1. Another way to illustrate the performance of a reinforcement learning agent is to plot the average total reward in each episode versus the number of experienced episodes, as in Figure 5.2. The first method is the more standard one used to compare the sample efficiency of different algorithms. However, plotting against the number of episodes should give equivalent conclusions, as better performance in one type of graph, meaning higher rewards in at least some sections of the graph, implies better performance in the other. It is also important to consider the fact that the learning process is random and the performance varies between trials hence the graphs were plotted using values averaged over multiple trials.

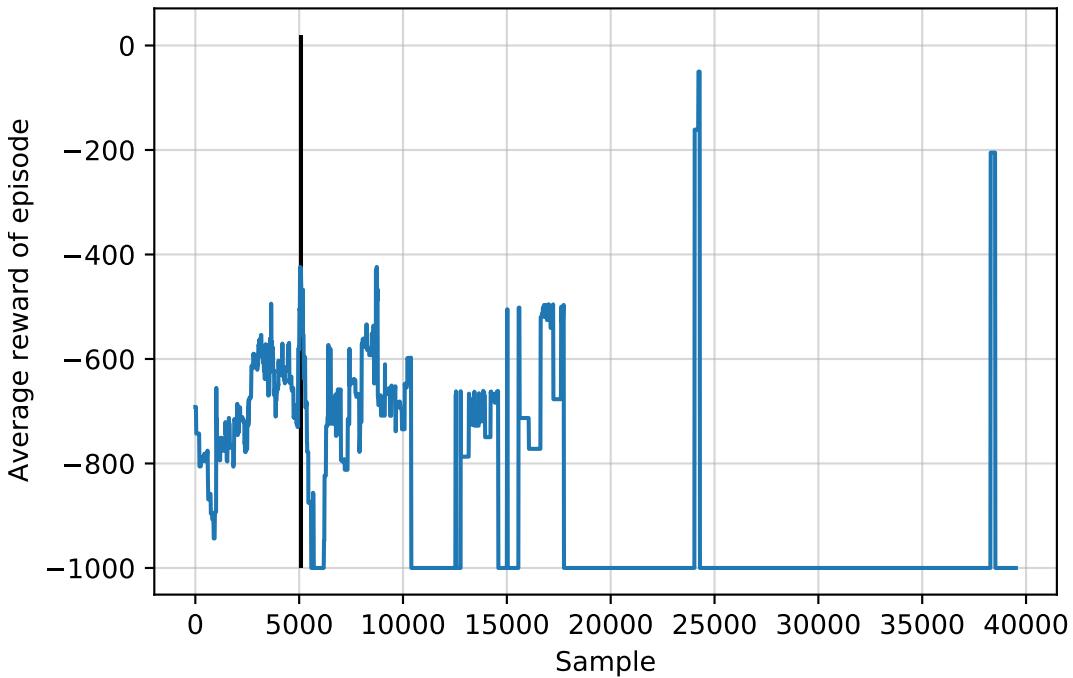


Figure 5.1: An example of an 'average reward' vs 'samples' graph. The red line indicates the length of the shortest trial. Moving to the right of the line, the quality of the graph decreases and the average reward counterintuitively decreases due to the bias of only "bad" trials having such a high sample count.

The main method of presenting the results in this report is using the 'total episode reward' vs 'episode' graphs. The reasoning for relying less on the 'total episode reward' vs 'sample' graphs is that the data was collected in experiments with a fixed number of episodes rather than a fixed number of samples. As the number of samples in each episode varies, each trial has a different number of samples. Thus, a graph of the average rewards can only be plotted for the length of the shortest trial. Plotting for any higher number of experienced would mean that the right section of the graph is an average calculated over fewer trials. This obviously increases its variance, but also introduces a bias. This is because a higher number of samples for the same number of episodes occurs when the episodes are longer, which is when the agent performs worse, as successfully completing a task early ends the episode. As a result, if the graph is plotted for more than the length in samples of the shortest trial, the rewards values for high sample counts are negatively biased and as some of the trials last a very short number of samples, the number of reliable samples is small in these graphs. For the experiments in this chapter, the 'total episode reward' vs 'sample' graphs, cut off at the length of the shortest trial, are provided in the Appendix.

5.2 Effect of different state representations on performance.

It is initially necessary to decide how to encode the states to be fed into the neural networks of the agents. As a first approach, the state is fed directly as an integer representing the index of the state. This is the most compact encoding, which is what would be used in standard Q-learning, in which

the state is used merely as the index of a look-up table rather than as the input of a neural network. Though extremely compact, this representation of the state is very hard to decode for the neural network and its structure does not directly represent aspects of the state¹, leading to convergence requiring too many experience time-steps. Thus, it is decided to encode the state in three alternative representations:

1. **Compact** The idea is to give a physical meaning to the state representation. The state is given by a vector of values, each representing a different aspect of the physical meaning of the state. In the 'Grid' environment, these are the row in which the Taxi is currently located, the column and its current destination. A lack of physical meaning of the state representation also would make it difficult for the neural network to generalise to unseen states, which is an important feature in large state spaces. Thus this representation could be the most useful in very large state spaces.
2. **One-hot encoding** This aims to make it easier for the neural network to learn to decode the state. The state is simply converted to a one-hot encoded version, which is a vector of length equal to the number of possible states, where all elements have value 0 except one which has value 1.
3. **Hybrid encoding** This combines the previous two methods. The destination is represented by a single bit, taking value 0 or 1, while the position of the state is represented by a one-hot-encoded vector. This results in a state representation that is a vector of size one bit greater than the number of possible positions in the environment.

¹As the state is represented by an integer, this has no physical meaning before decoding. However, the state could be represented for example by a vector of values, each having a physical meaning such as horizontal position, etc.

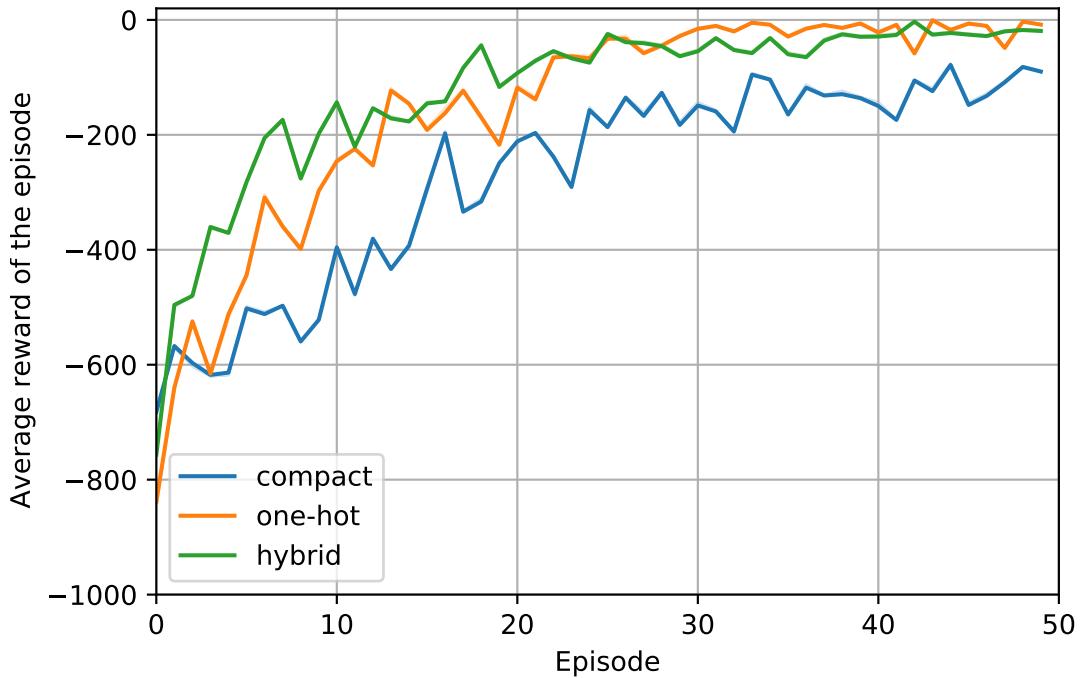


Figure 5.2: Effect of different state encoding methods on the performance of the baseline DQN agent in the 'Grid' environment.

Figure-5.2 shows that in the 'Grid' environment the hybrid encoding achieves the best performance, as the rewards are higher during the training phase, while they are similar upon convergence when compared to standard one-hot encoding. The compact encoding shows the worst performance at all points in time. It is decided to use the hybrid encoding for all experiments to expedite training times and hence experimentation times in this work. Long training times were a major issue in this project, thus increases in performance are welcome as they allow to run more trials to obtain higher quality results and to also run more experiments.

5.3 Effect of varying the size of neural network

The size of the neural network is another factor that influences the speed of learning of deep reinforcement learning algorithms. Figure 5.3 shows that for a DQN agent in the 'Grid' environment, the performance improves as the network size increases, at least for the tested range of dimensions. The 'small' network has roughly a fifth of the number of parameters that the 'medium' network has, while the 'big' network has roughly 25 times the number of parameters of the 'medium' network. When the network is too small, performance degrades as the network is not complex enough to represent the function that maps from states to action values of the environment. In the figure, the three bigger networks are all complex enough to reach convergence, however the biggest one does so quicker. This may be due to having more weight combinations giving the correct outputs for the possible states. While this leads to faster convergence and is not a problem in the environments used in this report, due to being able to visit all the states multiple times, it might lead to problems in

environments with very large state spaces that rely on generalisation to unseen states. This is because using bigger networks can cause overfitting, which hurts the generalisation properties of the network to unseen states. In this project, the neural networks are large enough to capture the dynamics of the environment, but not excessively large as this is what would be needed to avoid overfitting in more complex environments. The figure also shows that the difference in performance of 'medium2' and 'medium', which has half the number of trainable weights of 'medium2' is not very significant. This demonstrates that a change in number of parameters by a factor of 2 does not have a significant effect on the performance. In this project the total number of parameters is made approximately equal between different experiments. For example, when comparing an architecture having 1 NN with another having 2 NN, the networks of the latter would be smaller so that the total number of parameters would be approximately equal. This approach does not provide a mathematically fair comparison, however, as previously shown, small differences in number of parameters do not have a significant impact, thus it is unlikely that this affects the results enough to invalidate them.

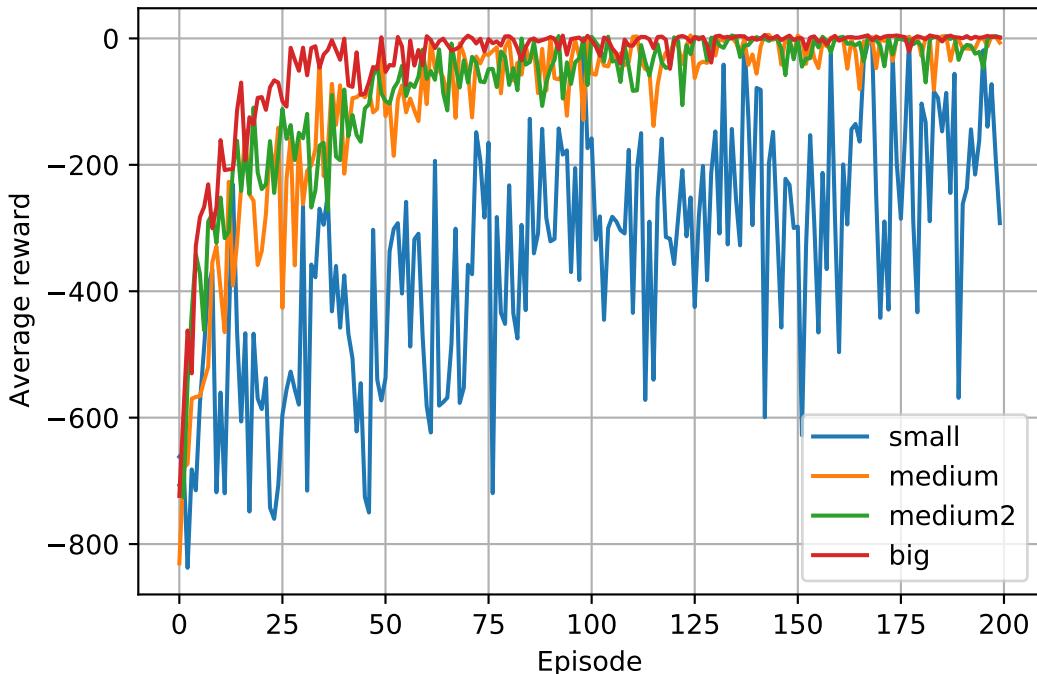


Figure 5.3: Performance of DQN with NNs of different sizes in the Taxi2 environment. 'medium2' has twice the number of trainable weights of 'medium'.

5.4 Initial results for 'Taxi2' and original state-decomposition architecture

The original state-decomposition architecture proposed in the previous chapter is initially tested in the 'Taxi2' environment. The results are shown in Figure 5.4, where the networks are combined at episode 100. In the first stage of the training, convergence is not achieved, as the transitions between subspaces are not considered. In the second part of training (starting at episode 100, when the

networks are combined), the average reward starts at the same levels as an untrained agent and the reward graph is not any better than that of a standard DQN agent. This means that the knowledge learned in the first part of training is 'forgotten' in the second part, thus not contributing to any faster learning in the second stage.

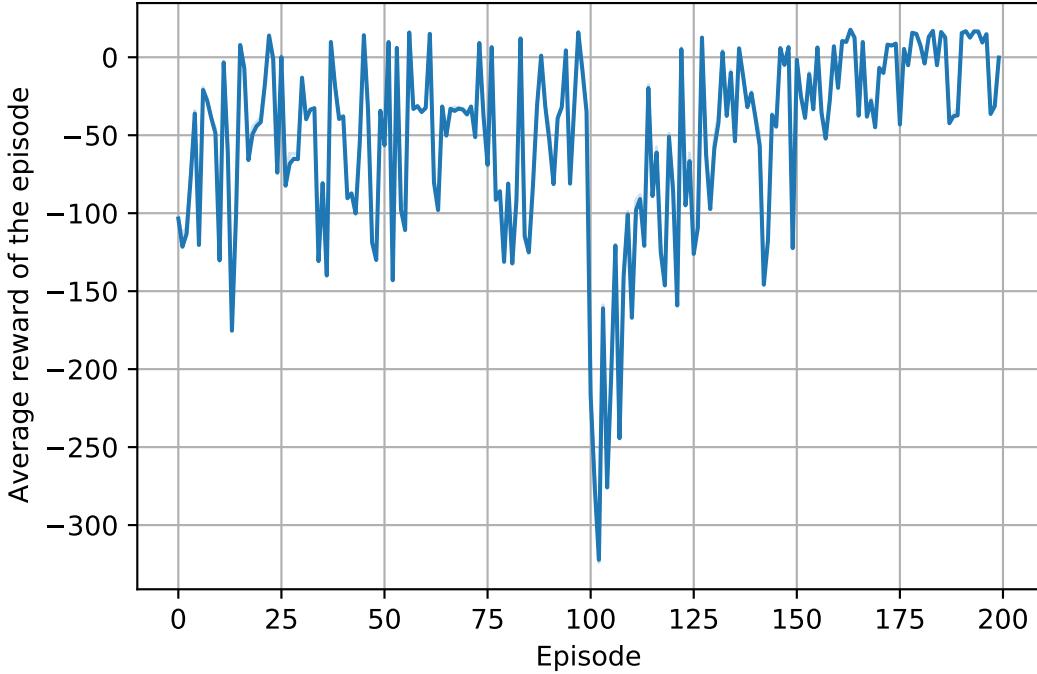


Figure 5.4: Performance of the original two-stage state-decomposition method. The networks are combined at episode 100, where the rewards fall to 'untrained agent' levels and take long to convergence, not any faster than a DQN agent.

5.5 A simple approach to decomposition: 'DeltaSwitch'

Following the failure to obtain positive results using the original method that involves the use of a combining neural network on top of the neural networks corresponding to each one of the subspaces, it is decided to experiment with a simpler approach in order to first validate the idea of state-decomposition in the simplest possible application. For this purpose, a network called 'DeltaSwitch' is introduced. This consists of one neural network corresponding to each subspace, as in the originally proposed architecture. The main difference is that the action values are combined using a simple weighed sum, so that the output action values of the "switch", as shown in Figure 5.5, are given by

$$q(s, a) = \sum_{i=1}^N \delta q_i(s, a) + (1 - N\delta) \cdot q_{\{j: s \in S_j\}}(s, a) \quad (5.1)$$

where $q_i(s, a)$ is the action value of action a given by the network of the subspace S_i with index i , δ is a small constant that satisfies $\delta \in [0, \frac{1}{N}]$ and N is the number of subspaces. This approach is also simpler in the fact that it only consists of one stage of training. This is done to avoid the "forgetting"

that is observed in the original method, which made the first stage of training effectively useless. In fact, all the networks are combined from the start in a predefined manner, as δ is a non-learned constant. The combining by weighed average effectively substitutes the combining neural network and acts as a switch when $\delta = 0$, as it effectively selects the output of the network corresponding to the subspace of the input space, thus the name 'DeltaSwitch'. This is the specific case shown in Figure 5.5. The idea of this approach is that it may not be necessary to combine the action values of the networks in a complicated manner and that doing so using another neural network might be complicated and lead to stability issues and lack of convergence. The updates of the network are carried out by treating the whole architecture as a single unit. δ is chosen to be small, so that networks other than the one corresponding to the input state only have a small impact on the output values. As a results, each sample used in an update of the network causes significant change in the weights of the network of the corresponding subspace, while only slightly changing those of other networks. This can be confirmed by deriving the formulas of the updates of weights given the loss function and optimiser, however it is sufficient to understand that those weights that have a smaller effect on the output are updated by a smaller amount. As the network is always treated as a whole, the target values required in the update of the network are generated using all the possible networks. For example, if a transition is from a state in subspace 0 to a state in subspace 1, the network that is updated most significantly by this sample is the one corresponding to subspace 0 while the target value is generated also using the outputs of the network of subspace 1. This means that each network is trained using target values generated by other networks, thus the networks are not independent during the training stage. This is required because 'DeltaSwitch' is a one stage method, thus the transitions between subspaces must be accounted for since the start of the training. Contrarily, the originally proposed algorithm could initially ignore transitions between subspaces as they were later accounted by using a combining NN.

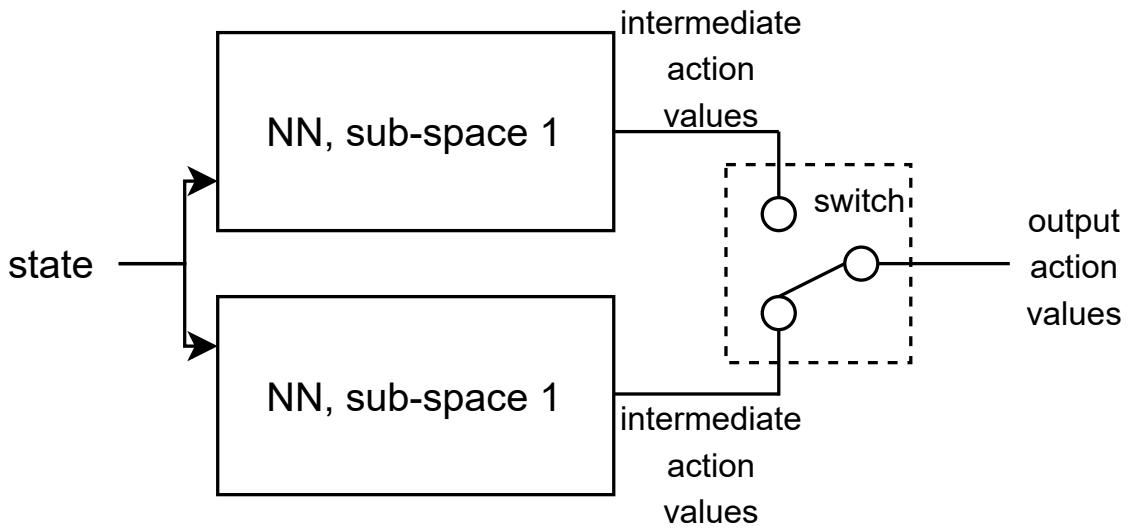


Figure 5.5: 'DeltaSwitch' architecture with $\delta = 0$. In this case, the weighted sum effectively acts as a switch, which selects the output of the network corresponding to the subspace of the input state.

The first environment where this is tested is the 'Taxi2' environment. The state space was decomposed into two subspaces, by assigning each state to the subspace corresponding to the closest

destination. In this environment there is no significant difference in performance between DQN and 'DeltaSwitch' for differing δ values, as shown in Figure 5.6.

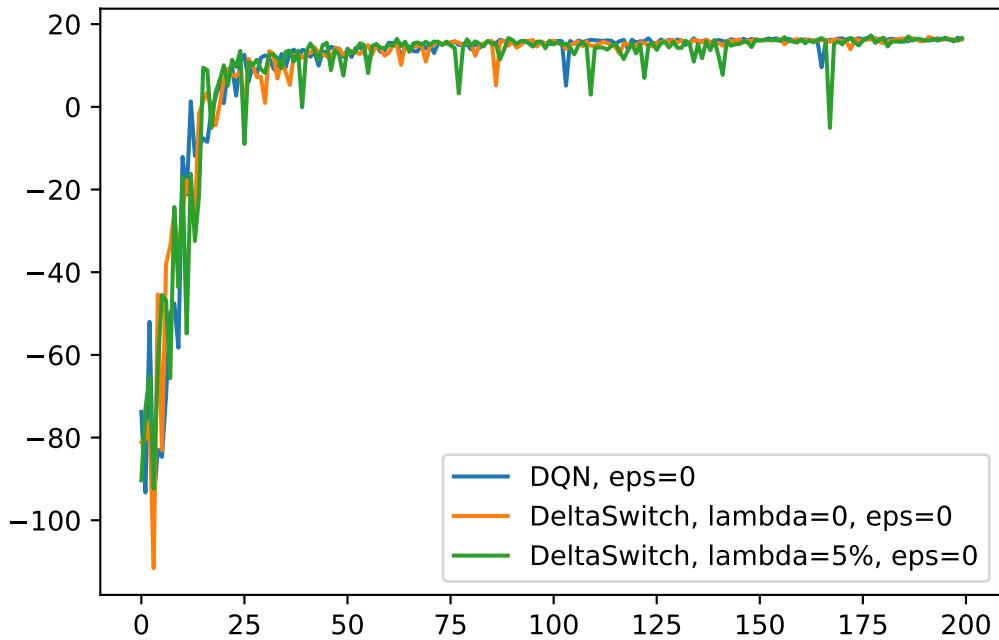


Figure 5.6: Comparison of DQN and DeltaSwitch with $\delta = 0$ in the 'Taxi2' environment.

While this does not show that the state-decomposition can lead to performance gain, it shows that it can at least have equivalent performance to a simple DQN agent. The experiment is then repeated in the 'Grid' environment for easier interpretability of the results. Figure 5.7 shows that here 'DeltaSwitch' achieves higher rewards compared to DQN when decomposing by destination. Decomposing by destination means that the possible states are divided in two subspaces, corresponding to the two destinations. Decomposing by position means that one subspace is formed by all those states in which the current position in the upper-left triangle of the Grid, while the other subspace is formed by those positions in the lower-right triangle.

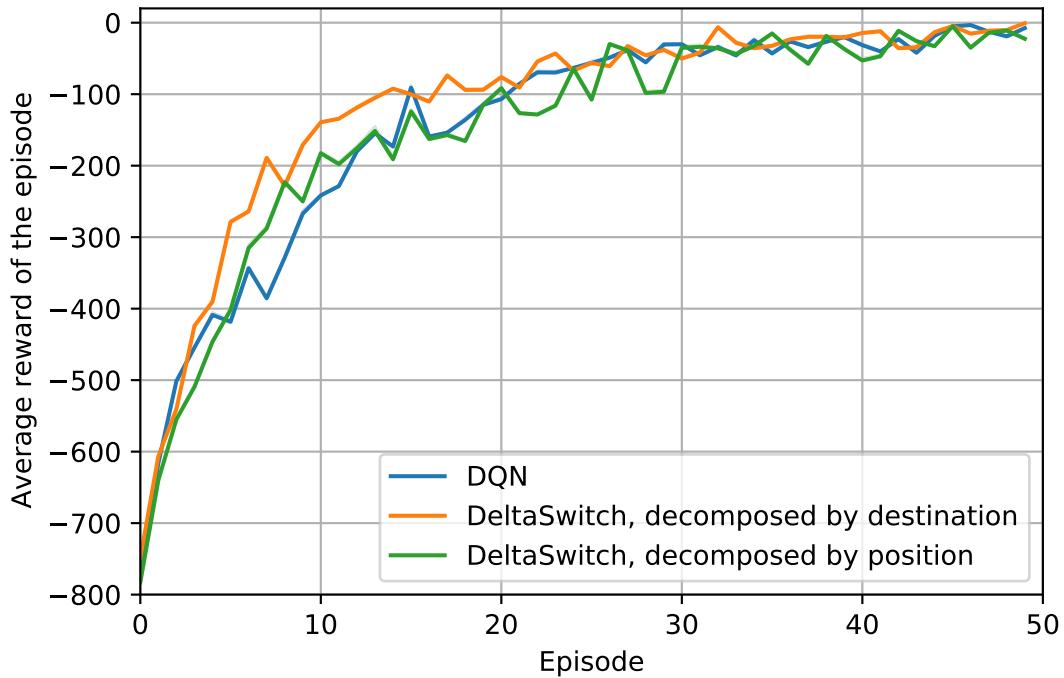


Figure 5.7: Comparison of DQN and DeltaSwitch with $\delta = 0$ in the Grid environment.

A possible reason for the performance gain is that the neural networks of 'DeltaSwitch' represent simpler functions when the decomposition is by destination. The main difficulty in the function representation is that states which have the same position but different destinations have very different action values. This difference is dependent on the position. Decomposing by destination allows each network to only deal with states with the same destination, thus greatly simplifying the function that the networks have to learn and leading to faster learning. The experiment has been repeated using standard one-hot-encoding and the results are shown in Figure 5.8. Here the differences between DQN and 'DeltaSwitch' decomposed by destination are even more marked, as QN and 'DeltaSwitch' with decomposition by position need to learn to decode the one-hot encoded vector to determine the destination, which is even harder than having the destination given as in the hybrid encoding. Additionally, 'DeltaSwitch' by position only has half of the training samples of DQN for each network to learn to do so, leading to a further decrease in performance.

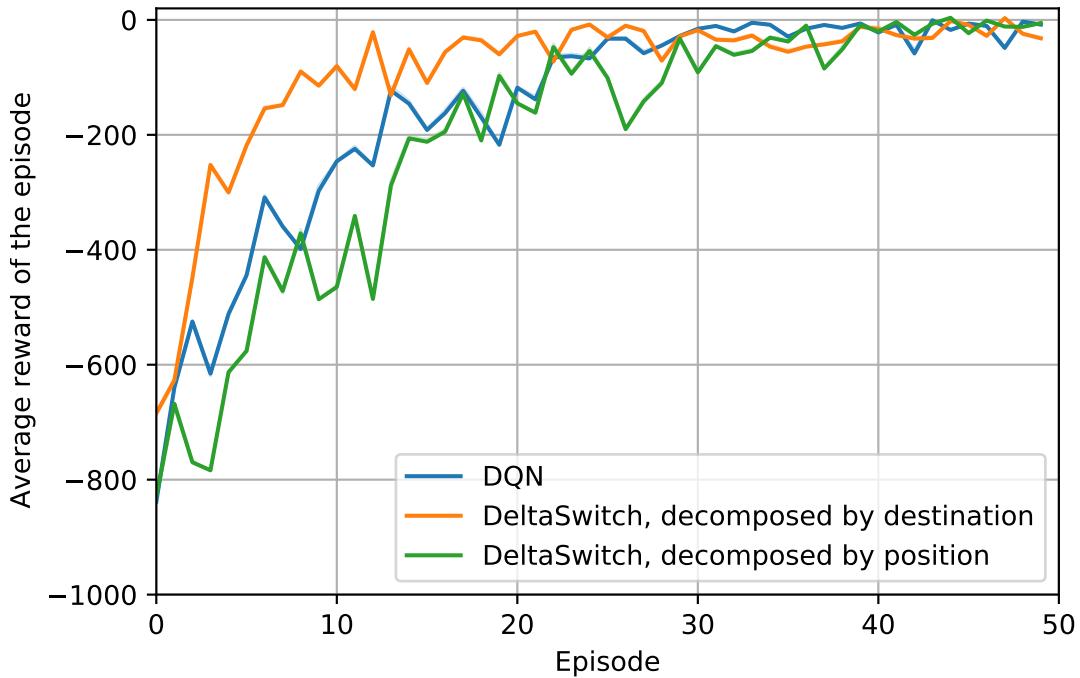


Figure 5.8: Performance of DQN and 'DeltaSwitch' when the state is one-hot encoded.

This result shows that 'DeltaSwitch' can present some performance gains compared to a normal DQN algorithm in a perfectly decomposable environment. This environment is perfectly decomposable by destination, as the destination never changes, meaning that there are no possible transitions between states having different destinations. It is also perfectly decomposable by position, as when adopting the optimal policy, the agent never moves from one of the triangles to the other, since the destination is always the one closest to the current position. This experiment validates the idea of the state-decomposition having a performance gain in such a specific environment and thus it may be possible to also have a performance gain in environments that are not perfectly decomposable.

For this kind of perfectly decomposable environment there is no reason to set $\delta \neq 0$ as the two subspaces are completely independent. Figure-5.9 shows how the performance of 'DeltaSwitch' degrades in the 'Grid' environment when increasing δ . In fact, there is never a reason to set $\delta \neq 0$, as even when the subspaces are not completely independent, as communication between them is allowed during training, the action values of the networks are accurate, so the weighted sum should act as a mere switch rather than combining multiple values. The idea of a weighted average has actually been devised to consider the case when communication between networks does not occur during training². However, we found out that this should be allowed, meaning that there is no reason to set $\delta > 0$.

²This refers to the generation of target values being allowed to use the networks of any subspace.

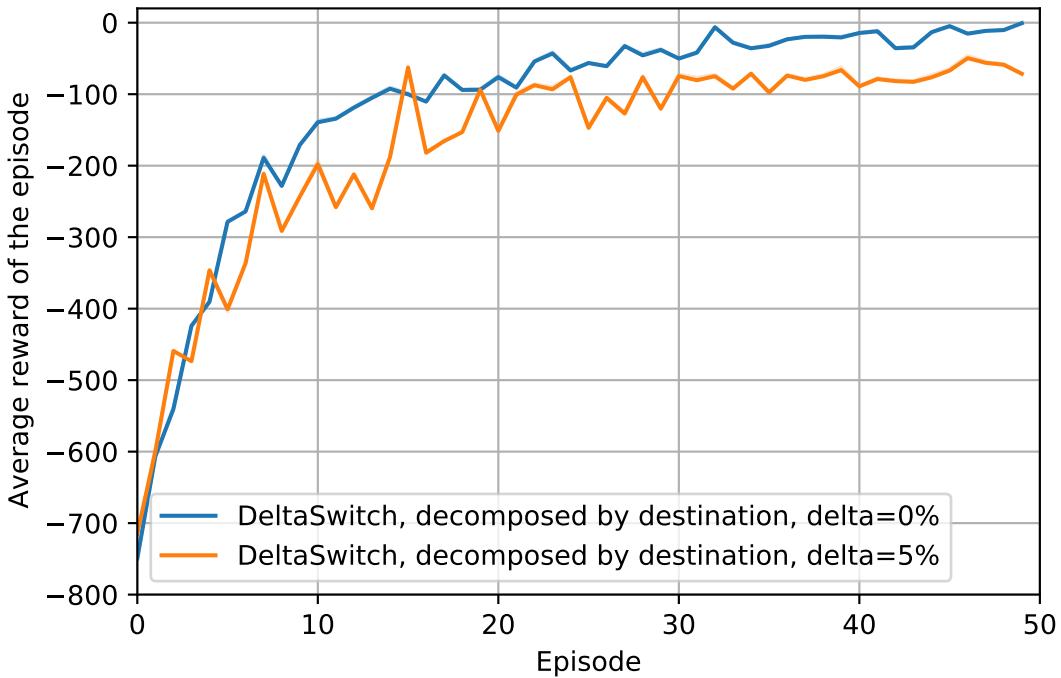


Figure 5.9: Comparison of DeltaSwitch with different δ .

5.5.1 Applying 'DeltaSwitch' to non-prefectly decomposable environment

After determining that a 'DeltaSwitch' agent with $\delta = 0$ can lead to faster training, we shall explore the performance of such an agent in environments that are not perfectly decomposable. This was done by running experiments in the 'GridEps' environment with $\epsilon > 0$. Figure 5.10 compares the performance of training a DQN agent and 'DeltaSwitch' agents, decomposing by destination and by position. It can be seen that for $\epsilon \leq 3\%$ 'DeltaSwitch' with decomposition by destination learns noticeably faster than the other two algorithms. The difference is especially large for $\epsilon = 3\%$, in Figure 5.10 (c). This might be due to this environment being the harder to solve, giving a greater margin of improvement over DQN. For higher ϵ values, there is no statistical difference between the three algorithms. This is probably due to the fact that the environment is not as decomposable anymore, as the transition probabilities between the subspaces are no longer negligible.

This is an important result, as it shows that while the state-decomposition method is applicable to perfectly decomposable environments, or environments that are sufficiently decomposable, the performance gain disappears once the transition probabilities between the subspaces are too large for decomposition.

Similar results were also obtained in the 'GridEps2' environment, which has slightly different state transition probabilities. In this case though, the decomposition by position was actually the most effective. It is not easy to reason why this is case, however it is an interesting result. This is because the decomposition by position is a different type of decomposition that does not rely on ignoring transitions with small probability in order to decompose the state space, but it simply decomposes

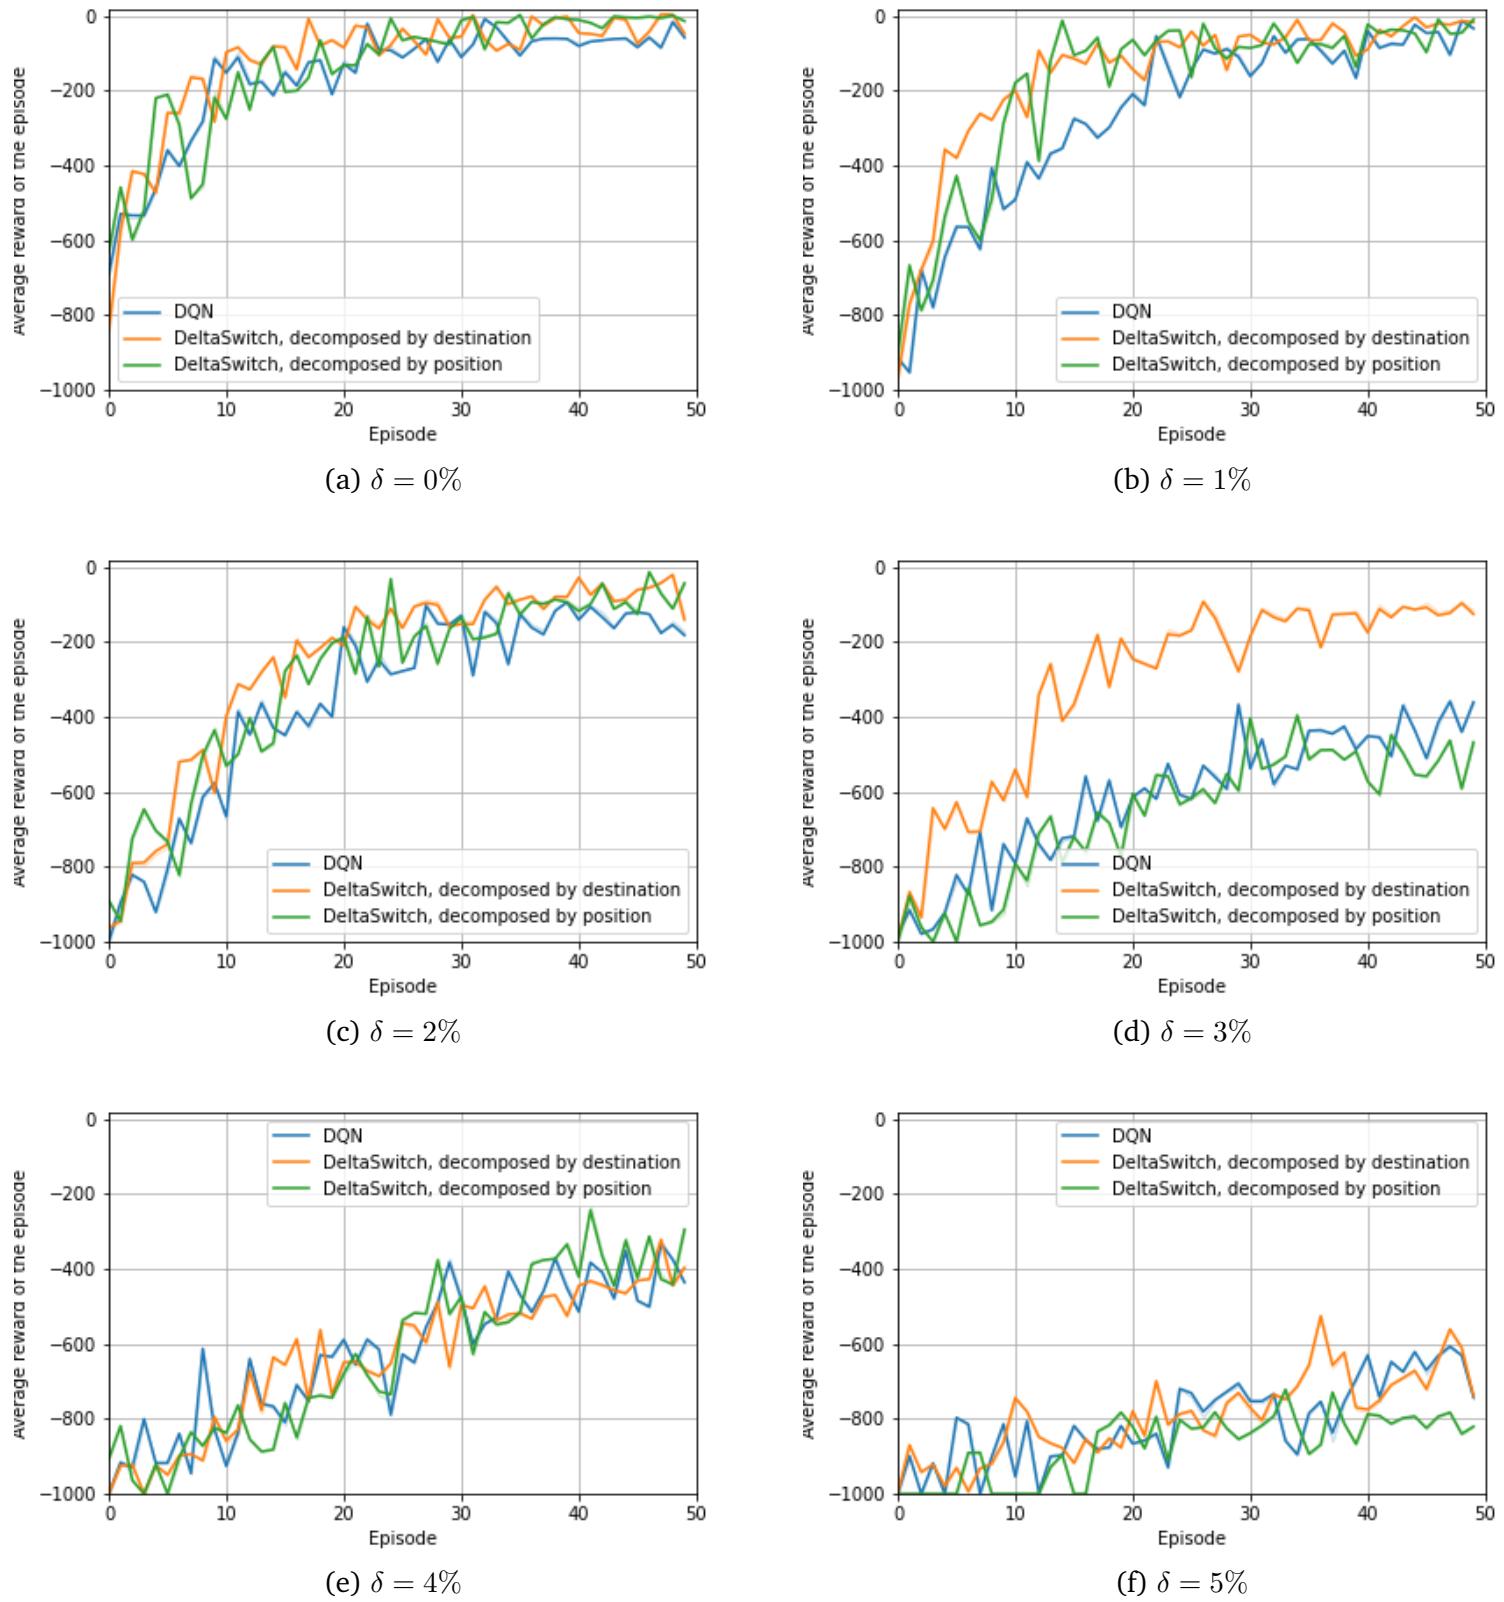


Figure 5.10: DeltaSwitch vs DQN in 'GridEps' with $0 \leq \epsilon \leq 5\%$.

the state space in order to divide the problem into two sub-problems, without considering state transition probabilities. This environment thus shows that in 'DeltaSwitch' the originally proposed way to decompose state spaces is not always necessarily the best. This is because as the networks are allowed to "communicate" during training as previously explained, we have the freedom to choose different types of decomposition, without having to limit ourselves to semi-independent subspaces. This result shows that the decomposition method can provide better performance even in cases with high transition probabilities between different subspaces.

At this point we can theorise that the mechanism by which 'DeltaSwitch' increases the efficiency of learning is by simplifying the functions that the networks need to learn. Considering the previous paragraph, we also conclude that the lack of performance gain of 'DeltaSwitch' over DQN when the transition probabilities between subspaces are large is likely to be due to the functions to be learned by the NNs of 'DeltaSwitch' no longer being much simpler than the function learned by the NN of the DQN agent, as is the case for small ϵ values. The following sub-section highlight some of the experiments that can be used to shed more light on the mechanism by which 'DeltaSwitch' achieves performance gain.

5.5.2 Effect of using reduced encoding

For the 'hybrid' encoding described in a previous subsection, the input to the first layer of the network is a vector of '0's and '1's of size the number of possible positions in the grid plus one. However, for $\delta = 0$, since each network only contributes to the output when the input state is in the network's subspace, it is possible to reduce the size of the input vector to only account for the states that are relevant to each network. For example, when decomposing 'Grid' by position, the two networks only deal with 312 and 313 positions instead of the original 625. This means that instead of the original vector of $625 + 1 = 626$ elements, it is sufficient to feed the two networks with vectors of $312 + 1 = 313$ and $313 + 1 = 314$ elements respectively. Experiments have been run to determine whether such a reduction in input size would improve the performance. Figure 5.11 shows that the performance gain when reducing the input size seems to be marginal and might be caused by the randomness of the trials.

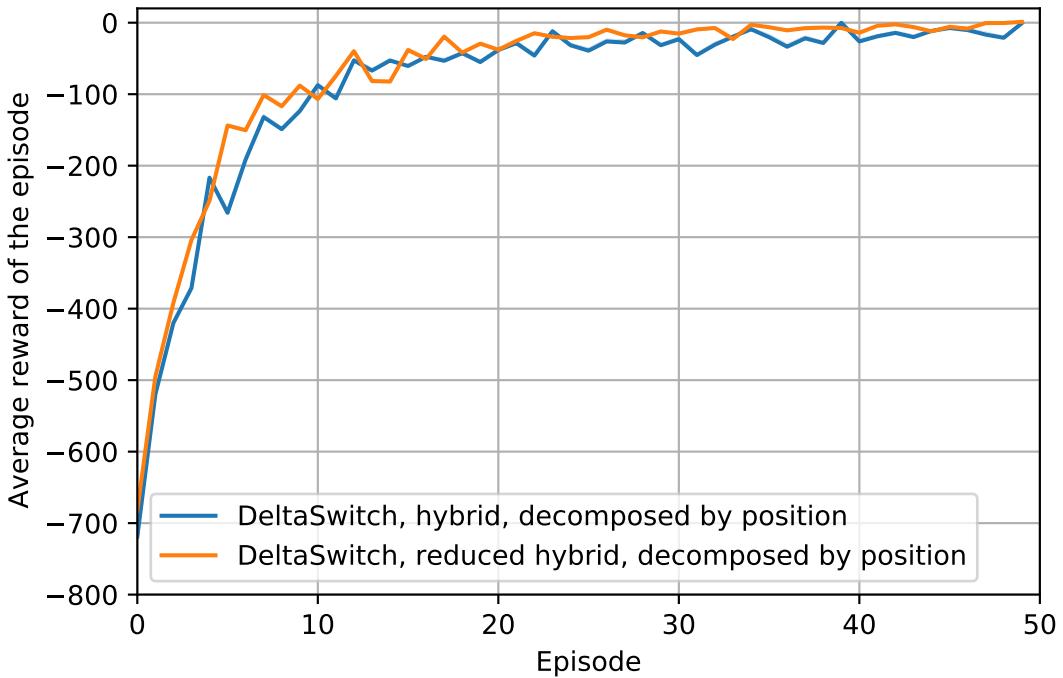


Figure 5.11: Effect of reducing the input size.

Effectively, when using the full input vector size, although the input vector is longer, almost half of its elements are always 0 when the network is active, meaning when the input state is in the network's subspace, so that the network affects the output³, thus these inputs mathematically have no effect on the network and it is intuitive that the results should be similar when using the reduced input vector.

5.5.3 Effect of adding a wall to separate subspaces

The experiments with 'DeltaSwitch' are repeated in a modified version of the 'Grid' environment. This version has a wall that separates the upper-left triangle of the grid from the other triangle. This is done to test the effect of having decomposition upon convergence, versus having decomposition with any possible policy. When there is no wall, the decomposition by position only occurs upon convergence, that is when the optimal policy is being used, as there is no reason why the agent would act to cross the boundary between the two subspaces (i.e., positions in the upper-left and the lower-right "triangles"). However, it is physically possible to cross this boundary and this is very likely to occur during the training stage. Conversely, when a wall is present, it is not physically possible to cross this boundary, regardless of the policy.

It is intuitive that when the wall is present, the learning process is faster, as it impedes the agent from moving further from its destination. It is interesting to notice that as shown in Figure 5.12, the performance of 'DeltaSwitch' and the benchmark DQN algorithm is equivalent in this environment. Any differences in the figure are small enough to be attributed to the randomness of the training, as they are not consistent throughout the graph, with the lines crossing each other multiple times.

³Note that $\delta = 0$, so for other state the network does not affect the output and can be referred to as inactive.

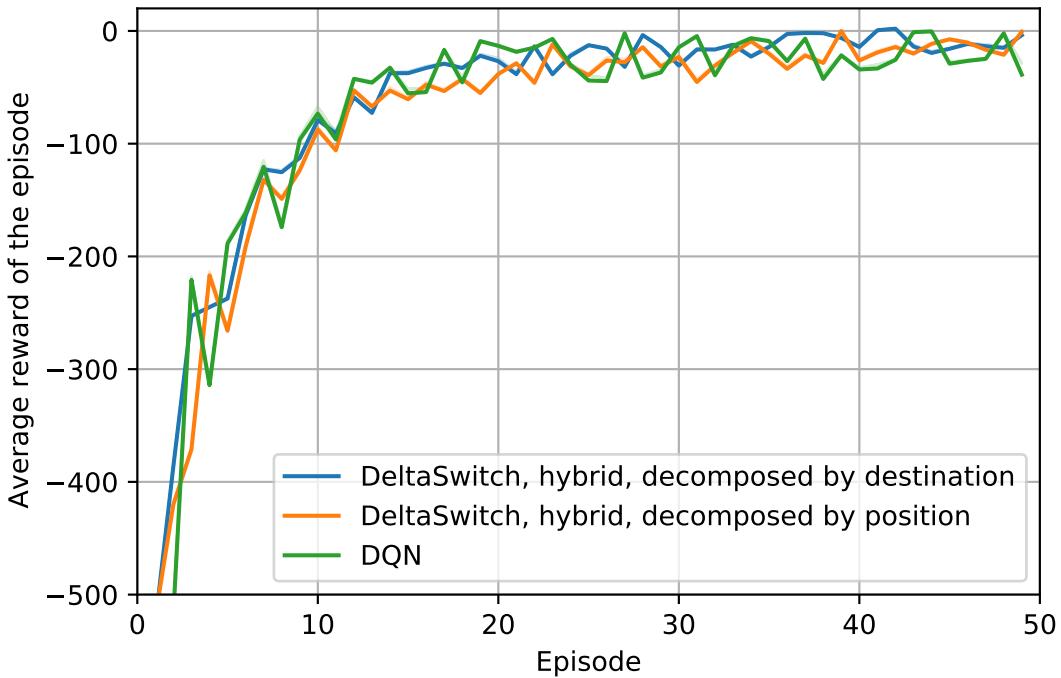


Figure 5.12: Performance of DQN and DeltaSwitch with different decomposition in the 'Grid' environment with wall.

Whereas in the environment without a wall, 'DeltaSwitch' with decomposition by destination presents a performance gain, this does not occur when there is a wall. This agrees with the theory that the performance gain of 'DeltaSwitch' with decomposition by destination is due to the networks having to learn simpler target functions. When the wall is present, states in which the destination is far, that is, states that are on the wrong side of the wall, are never visited. As a result, for each position on the grid there is only one possible destination, meaning that the function to be learned by the network is already simplified, even for DQN and that applying the decomposition does not simplify it further, and thus the performance is very similar for the three methods.

5.6 A different approach based on transfer learning

The original idea of state-decomposition based on small transition probabilities exploits the fact that the optimal policy to be followed in a perfectly decomposable environment should be similar to that of a similar environment, having a small probability of transitions between states of the different subspaces. Another class of techniques that leverages similarity between environments is transfer learning. Transfer learning is a technique used in machine learning to transfer knowledge between different tasks, in order to more efficiently learn in new tasks. Different methodologies have been explored for classification, regression and clustering [20]. Using transfer learning in reinforcement learning means to use the knowledge obtained by having trained an RL agent in a certain environment in order to more efficiently learn in a new similar environment, as in [30]. Two-stage state-decomposition architectures learn local agents for the separate subspaces and a combining neural

network that together with the local agents, forms a global agent. This idea is very similar to training an agent in the perfectly decomposable environment and then using its knowledge, more efficiently learn in the more complex environment. Effectively this two stage approach is transfer learning from a simpler fully decomposable environment. Thus, exploring transfer learning techniques could give insight about possible ways to apply state-decomposition using a two-stage architecture, as shown in Figure 5.13.

It is shown in the previous sections that 'DeltaSwitch' can achieve performance gain with certain kinds of state-decompositions, however there are also different ways to decompose a state space that would require the use of a combining network. For example, consider having a state space represented by a $2n$ -dimensional vector, composed of $2 n$ -dimensional vectors. The dynamics of the model are such that transitions in which both sub-vectors change are very rare. Using the original decomposition method, this would not decompose. However, we could for example decompose by simply splitting the vector into the two sub-vectors, so that each of two neural networks deals with changes in one of the sub-vectors and a combining network deals with merging the results of these two networks. A practical example of such a situation is provided in the next chapter. Therefore, there are situations where a two-stage architecture involving a combining network would be required in order to apply the state-decomposition.

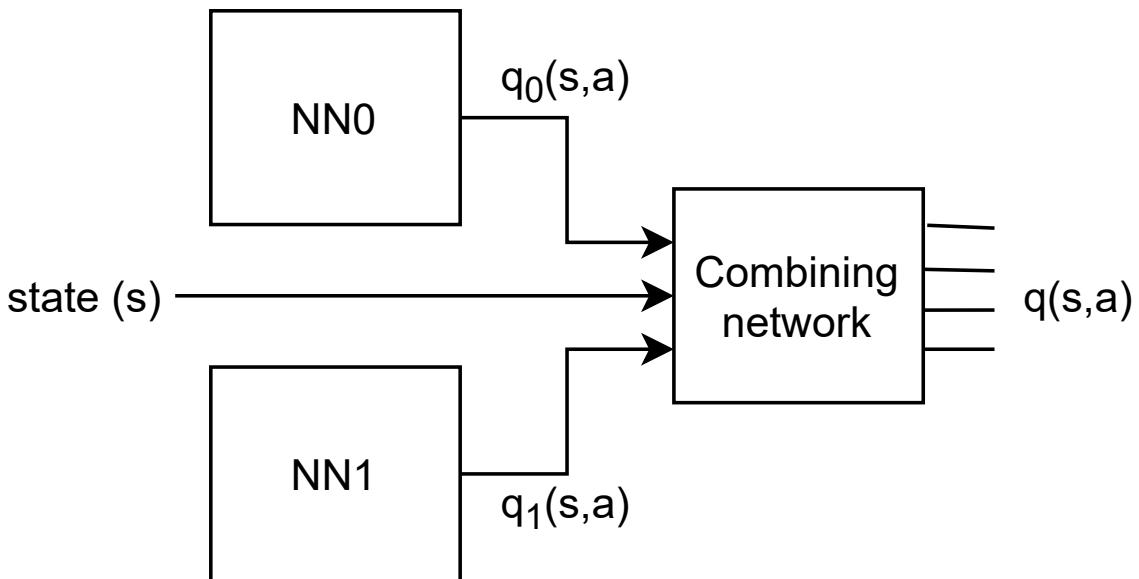


Figure 5.13: Combining 'NN0' and 'NN1' using another NN.

The following methods are tested by using the two networks from a 'DeltaSwitch' agent that has been trained in a 'Grid' environment, decomposed by position⁴, which gives two neural networks, 'NN0' and 'NN1', which correspond to the two subspaces. The transfer learning techniques are then tested in a 'GridEps2' environment.

The first approach is to combine 'NN0' and 'NN1' using another neural network, which replaced the "switch" in 'DeltaSwitch'. The weights of 'NN0' and 'NN1' are set to be fixed. This is basically

⁴This kind of decomposition was chosen as it has the best performance in the 'GridEps2' environment.

the architecture proposed in the original state-decomposition method, with the combining network also being fed the input state. The idea is that this network can act as a clever switch instead of the 'rudimentary' switch used in 'DeltaSwitch', in order to take into account the transition between subspaces. This algorithm gives the results in Figure-5.14, which shows that this algorithm, despite using 'NNO' and 'NN1' which have been previously trained, achieves lower rewards than a standard DQN algorithm for both values of ϵ . This gap increases for higher values of ϵ . This means that this method is unsuccessful in leveraging the knowledge obtained from 'NNO' and 'NN1'

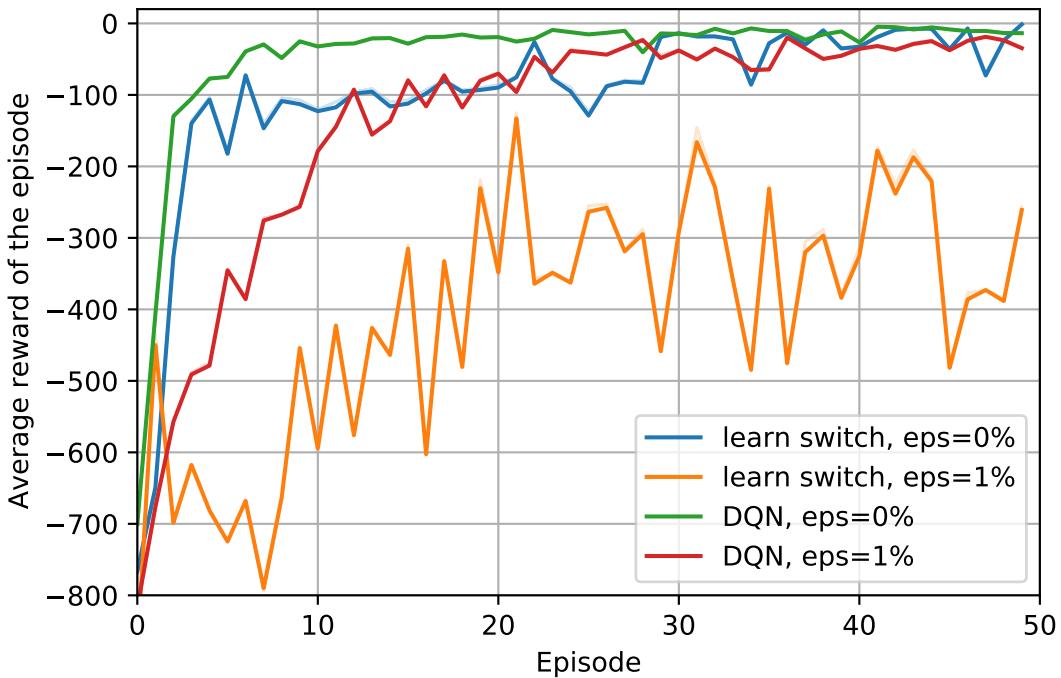


Figure 5.14: Performance of combining 'NNO' and 'NN1' using another NN that is also fed the current state as input in a 'GridEps' environment with $\epsilon = 0\%, 1\%$.

As the ϵ of the 'GridEps2' environment is set to be small, it is fair to assume that directly using the original 'DeltaSwitch' agent that 'NNO' and 'NN1' are taken from should give near-optimal performance. Thus, at least in the start, the transfer learning method should achieve much higher rewards than a standard DQN. This can be achieved by "forcing" the initialisation of the combining network, whereas before this was randomly initialised. This was done by training the combining NN in a supervised way so that the combined network ('NNO', 'NN1' and combining network) is initialised to behave exactly as the 'DeltaSwitch' agent that 'NNO' and 'NN1' are taken from. The training dataset for the "forced initialisation" has as inputs all the possible states and as outputs the corresponding outputs of 'DeltaSwitch'. After training for multiple epochs⁵, with MSE as the loss function, the combined network can be utilised in the environment as before and the RL process can be started. The initial exploration rate is set to be low⁶ as it is assumed that it is not needed to perform lots of exploration,

⁵In the experiment, the network was trained for 1000 epochs, in order to ensure an accurate initialisation. However, such high accuracy might not be needed.

⁶The ϵ of ϵ -greedy

due to the good initialisation. Figure-5.15 shows the results of this methodology with different NN sizes. The performance is much better than with random initialisation of the combining network.

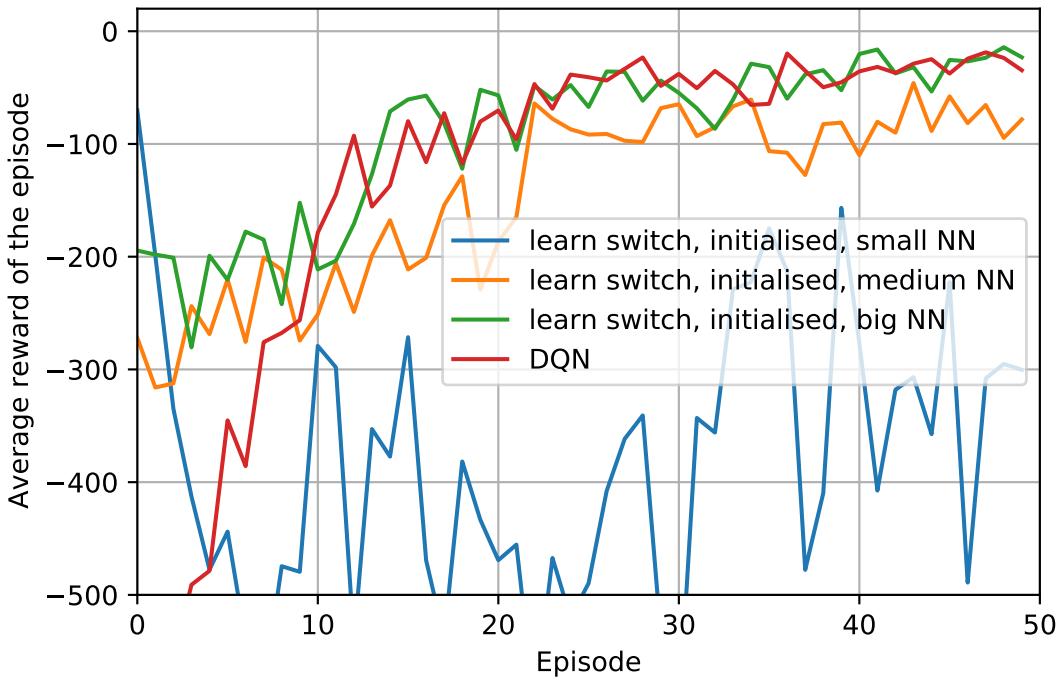


Figure 5.15: Performance of combining 'NNO' and 'NN1' with another NN when initialising the combined network to behave as the original 'DeltaSwitch'. Tested with different network sizes

Using the smallest network does not allow to capture the environment's dynamics and fails to converge. Only the biggest NN allows to achieve performance comparable to DQN in the last part of training and in this case the overall network size used is actually bigger than that in DQN. Requiring a combining network of a similar size to that of the DQN agent means that the function that this NN has to represent is similarly difficult to that of the DQN's network. Thus it seems that the initialisation of the network only serves the purpose of limiting negative rewards at the start of training, but it does not have a lasting effect in making the learning easier. The issue could be that, though Q-learning is guaranteed to converge after sufficient amounts of exploration, it always finds the optimal policy so it is likely to completely discard the initial policy provided when initialising the combined network, causing so called "forgetting". Therefore, SARSA is experimented and its updates are given by the formula

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]. \quad (5.2)$$

That is, SARSA is an on-policy algorithm, meaning that it updates its action values based on followed policy, rather than using the maximum action value in the next state as in Q-learning. The intuition was that doing so is more likely to preserve the initial policy provided by the initialisation of the combined network, which would increase the speed of learning. However, using SARSA failed to convergence, thus expected SARSA is experimented, which has the updates

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha\{R_{t+1} + \gamma \mathbb{E}_\pi[Q(S_{t+1}, A_{t+1})] - Q(S_t, A_t)\} \quad (5.3)$$

where π is the policy at the time of the update. Figure-5.16 shows that using expected SARSA seems to improve the initial performance, as rewards never fall below -200. However, in the long term the performance is still equivalent to DQN, as this method does not reach convergence earlier than DQN, thus the problem of "forgetting" is not solved.

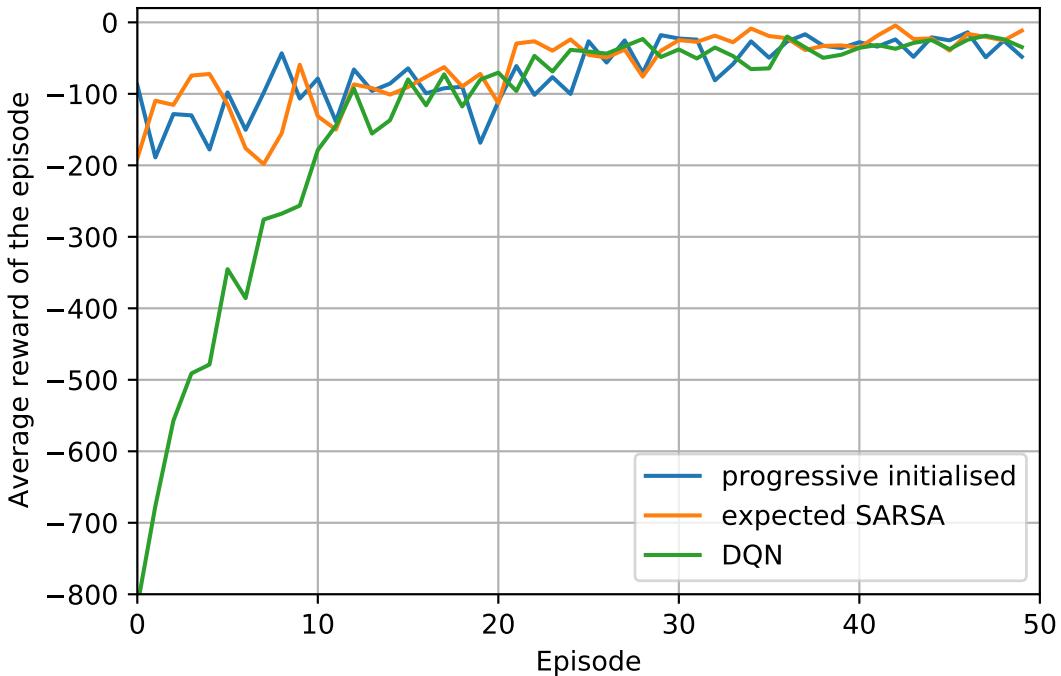


Figure 5.16: Comparison of using Expected SARSA after combining, progressive combined network and standard DQN.

Subsequently, the idea of 'progressive neural networks' [24] is studied. This is a kind of neural network that is devised for transfer learning, as it shows better performance in transferring knowledge between tasks and it avoids "forgetting". The progressive neural network is conceived for learning successive related tasks. In the first task, a neural network is used as in normal DQN. Then, in the second task, the neural network is formed by the neural network trained in the first task with its weights frozen⁷ and a new "column" of neural network layers. Each layer in the new column receives in input the outputs of the previous activation layer of the same column and the outputs of the activation layer of the network of the previous task. This combined network is then trained in the new task and the process is repeated for all the tasks, each new column receiving in input the outputs of all the previously trained networks. Figure 5.17 shows the structure of a progressive network used on a third task when transferring knowledge from two previously trained tasks.

⁷By frozen it is meant that the weights are not trainable.

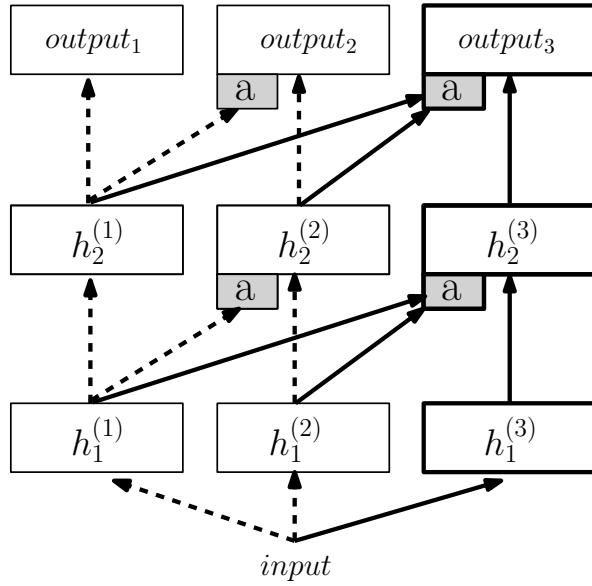


Figure 5.17: Structure of progressive network when training the third task. Figure taken from [24].

The architecture to be proposed is very similar. Instead of connecting the networks one after the other, a new network is connected to all the previously trained networks that correspond to the decomposed subspaces. That is, 'NNO' and 'NN1' have their weights and the outputs of all their layers connected to the corresponding layer inputs of a new network. This new network has thus the same number of layers as 'NNO' and 'NN1'. The architecture is shown in Figure 5.18. This is initially tested with random initialisation of the combined network, leading to slow convergence and worse performance than DQN. However, when initialising the combined network similarly to what is previously done in the other combining methodology, the performance is similar to what is achieved in the other methodology using expected SARSA, as shown in Figure 5.16. Thus, this technique successfully improves rewards in the initial training but also fails to provide an important transfer of knowledge between tasks and solve "forgetting", meaning that using the pre-trained 'NNO' and 'NN1' networks does not cause convergence to be achieved with fewer experienced samples.

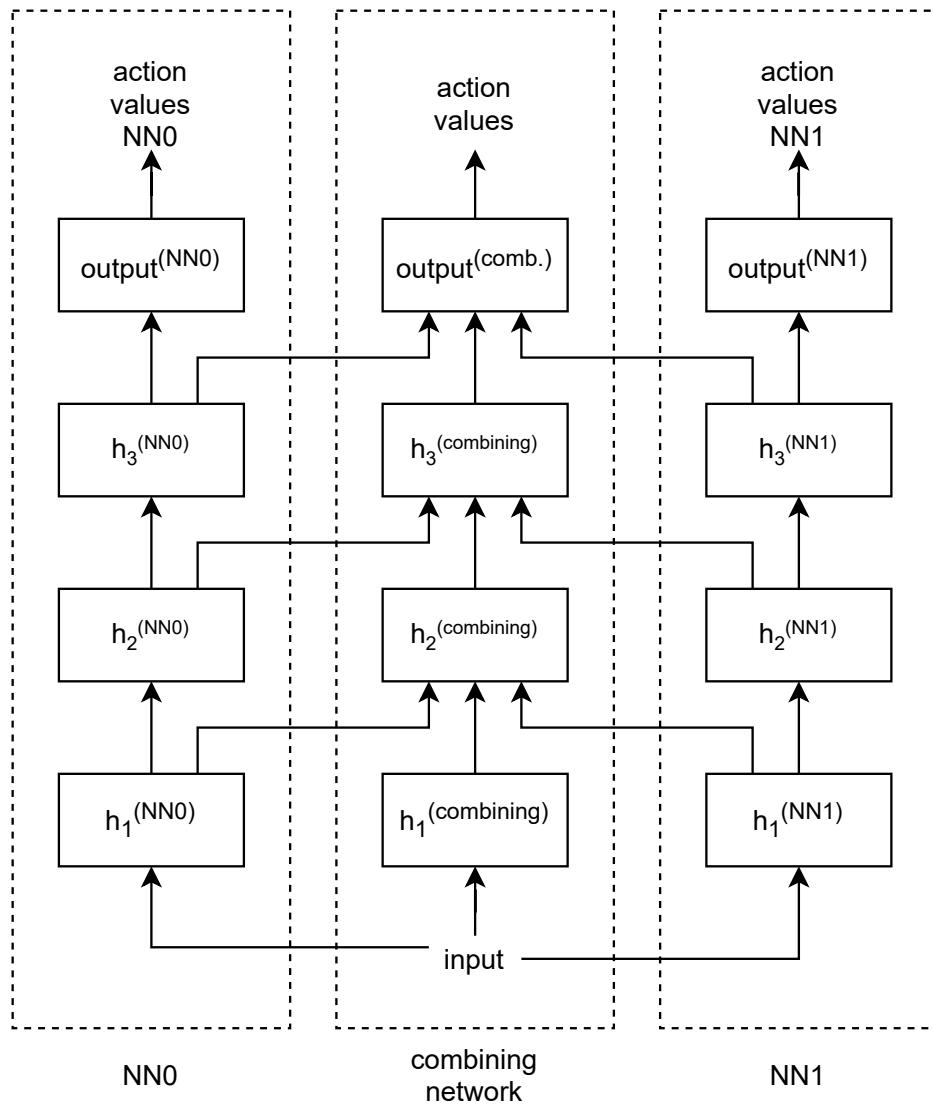


Figure 5.18: Modified progressive neural network. The layers of the combining neural network receive as inputs the respective outputs of the previous layers of frozen 'NN0', frozen 'NN1' and the combining network itself. h indicates hidden layers of the network.

These transfer learning techniques could be used to build a two-stage state-decomposition architecture in which in the first stage we consider the transitions within subspaces and in the second stage we apply transfer learning to learn the full environment, including all transitions. This was attempted without success, as providing a good initialisation for the second stage requires an almost optimal policy in the first stage, which takes many episodes to reach, thus not providing any benefit over DQN.

5.7 Techniques to learn the 'switch'

While the transfer learning techniques in the previous subsection show superior performance to DQN, this is an unfair comparison, as they rely on the previously trained networks 'NN0' and 'NN1'. Moreover, the better performance is due to a good initialisation which does not translate to a suc-

cessful transfer of knowledge and earlier convergence. However, the two combining techniques can be modified to work in one stage, thus simply providing a state-decomposition technique rather than doing transfer learning. This provides a more flexible architecture than 'DeltaSwitch' that could be used even when the networks of the subspaces⁸ do not account for all possible states, as could occur when using different types of state-decompositions as described in the next chapter.

The first proposed technique is to use a combining neural network right from the start, instead of training 'NN0' and 'NN1' first using a 'DeltaSwitch' agent. 'NN0' and 'NN1' are forced to learn the action values corresponding to the decomposed subspaces by including their outputs in the combined network's output, so that they are also considered in the loss function. Understanding this requires some detail on how target values are normally generated. In the standard DQN algorithm, at each update step, a batch of samples is selected from the replay memory. Each sample provides information about how the action value of a state-action pair should be updated, which can be used to set the target value of the state-action pair⁹. The target value of those state-action pairs that are not in any of the batch's samples are simply set to be equal to current output of the network for the corresponding states. This way, the network tends to improve for the state-action pairs in the batch's samples, while penalising changes in other pairs. A similar method is used to deal with the fact that only certain samples should update 'NN0' and the others should update 'NN1'. For those transitions samples that start in a state that is in subspace 0, the target value for the output of 'NN0' is set as normal using the formula of Q-learning updates, while the target value of 'NN1' is simply set to be what 'NN1' currently outputs for the considered state. The same is true the other way around. The output of the combining network is treated the same way as the output of the network in DQN, meaning that all samples are used to update it. This methodology allows the combining network to learn the true action values for all states, while 'NN0' and 'NN1' learn them for the states in the assigned subspace. However, as seen in the previous subsection, the combining neural network architecture does not seem to be able to make use of the information provided by 'NN0' and 'NN1'. As here 'NN0' and 'NN1' are not pre-trained, one can only expect the results to be worse. This is confirmed by Figure ??¹⁰, which shows that in 'GridEps2' with $\epsilon = 1\%$, this architecture performs worse than standard DQN.

The second proposed technique is based on the modified progressive network from the previous subsection. 'NN0' and 'NN1' instead of being pre-trained and having frozen weights are trained in a similar fashion as what was done in the technique in the previous paragraph. That is, the progressive network is composed of three columns, 'NN0', 'NN1' and the combining network which at every layer also receives the outputs of the previous layer of 'NN0' and 'NN1'. This is essentially, the same as the architecture shown by Figure 5.18, but 'NN0' and 'NN1' are now not frozen and not-pretrained. The only difference between this method and the method in the previous paragraph is given by the connections between 'NN0', 'NN1' and the combining network, but the way the target values are generated is completely equivalent. This second technique achieves better results than the previous one, probably due to the greater number of connections between the networks, acting at all the layers, which allows for a better flow of information from 'NN0' and 'NN1' and to share intermediate data-representations. Figure ?? shows that the performance is better than DQN and slightly worse than 'DeltaSwitch'. This means that this technique is capable of leveraging the state-decomposition

⁸In the examples, 'NN0' and 'NN1'.

⁹The target value refers to the output values used to calculate the loss function to run the update step of the neural network.

¹⁰All the experiments in this figure were run by applying decomposition by position.

in order to speed up the training. In the technique of the previous paragraph, the only connection between 'NN0', 'NN1' and the combining network was given by the outputs of 'NN0' and 'NN1'. This makes it easy for the combining network to completely ignore 'NN0' and 'NN1', thus making the training of 'NN0' and 'NN1' a useless complication that only harms the training performance.

Chapter 6

Applicability of the decomposition method

The previous chapter shows that given certain conditions, the state-decomposition method can give a performance gain in terms of the rewards obtained during the learning process. This chapter aims to provide a practical example where state-decomposition could be applied to a reinforcement learning solution and could achieve superior performance compared to a single global agent such as a standard DQN agent.

In general, techniques based on the state-decomposition method could be useful for any kind of RL problem in which using multiple function approximators to estimate the action-values allows to learn simpler functions than using a single approximator. This is often the case in systems composed of multiple sub-systems, where the various components seldom interact with each other, which is the case in which the originally proposed state-decomposition technique can be used. It is useful to discuss a possible application of the proposed methodology to better show how such a technique could be applied to a real-life scenario.

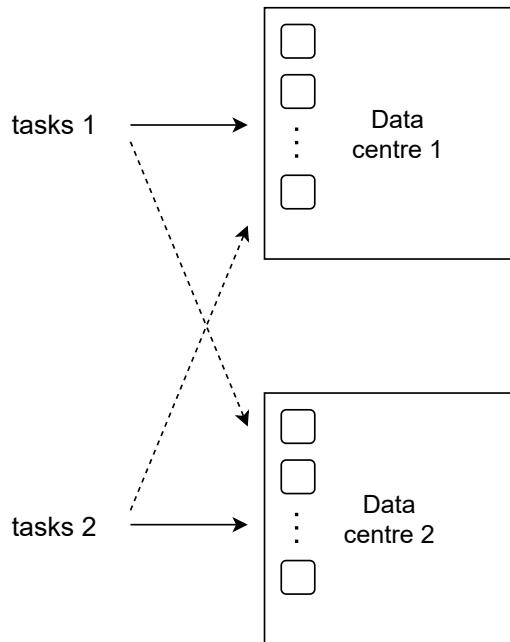


Figure 6.1: Allocation of resources problem for two interconnected data centres.

Consider a network scenario with two data centres, each of which has multiple computers, as shown in Figure 6.1. Each of these data centres receives processing tasks (jobs) that can be assigned to the computers and completed. Each of these jobs requires a random time to process on its allocated machine. At each time-step there can be no task or one task arriving at one of the data centres. The resource allocation problem deals with assigning these tasks to machines in the two data centres in order to minimise the average waiting time to completion of the tasks. Most of the time it is most efficient for the data-centre that receives the task to process it locally, as sending the task to the other data-centres incurs in additional delays. However, sometimes a data-centre is too busy to justify the additional delay incurred in transferring tasks between data-centres. This problem can be formulated as a Markov Decision Process and solved using reinforcement learning. The MDP can be defined as follows:

- The state of the MDP indicates the number of tasks waiting for each machine (a computer in a data centre). It also indicates the presence of eventual incoming tasks in the data centres. Thus the state can be represented by a vector having an element for each machine, which takes integer values indicating the number of jobs waiting for processing in the corresponding machine. The vector would also include two additional elements, taking values 0 or 1, indicating the presence of new incoming tasks in each of the data centres.
- The actions to be taken by the controller are to assign the tasks to the machines. At each time-step, the controller can allocate the received task to any machine, either in the local or in the remote data centre.
- The rewards can be positive values that occur when tasks are completed. In order to encourage local allocation of tasks, the rewards can be higher when task are completed in the data centre that receives them than when they are processed remotely.

If the difference of the rewards upon completion of tasks between local and remote data centres is set to be very high, the controller will allocate most tasks locally. The fact that the controller is unlikely

to assign tasks to the remote data centre makes it so that the equivalent transitions are unlikely to happen and that the state space can be decomposed accordingly.

The simplest way of applying state-decomposition in this scenario would be with a 'DeltaSwitch' agent having two neural networks, one for when a task is received by one data centre and another NN for when a task is received by the other data centre. Though this might be effective, it does not use the decomposition previously discussed based on decomposing by 'ignoring' state transitions below a certain threshold. In this case, an intuitive decomposition of treating two relatively separate problems as such is used.

In the case where the difference in rewards between completing a task locally or remotely is infinite, the tasks would always be solved locally. In this case the problem of deciding which data centre the task should be processed does not exist. In this case we could simply have two separate agents for the two data centres that learn how to assign tasks within the data centres, without dealing with their interconnection. However, as the difference in rewards decreases, it becomes optimal to use the interconnection. The 'learn the switch' and techniques discussed in the previous chapter would be useful in this case. The networks trained to assign tasks internally within the two data-centres can be combined in order to learn a global agent that can assign tasks in both in the most optimal manner. This architecture is illustrated in Figure 6.2.

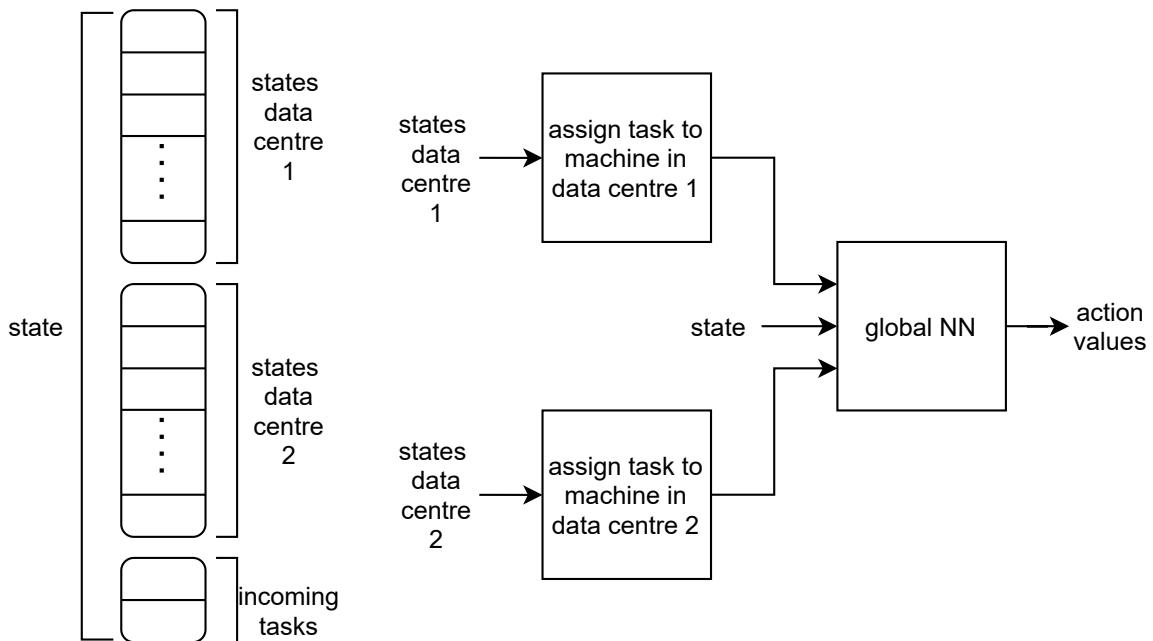


Figure 6.2: Possible state-decomposition agent.

This last methodology is very similar to what would be done in a hierarchical architecture. This would have an agent that explicitly decides whether a task should be assigned to data centre 1 or 2. Given this knowledge a local agent trained to decide how to assign the task internally within the data centre would be used and the outputs of these two would be combined to give the full action description. The difference with the state-decomposition method is that while in the latter the full action description is given by the combination network and all other components merely provide

implicit information to it, in the hierarchical approach the action description is formed explicitly by combining the outputs of multiple 'layers' of local agents. Additionally, the hierarchical approach deals with the problem with a 'high level' to 'low level' approach, in which 'high level' actions are decided first and 'low level' actions are determined within the 'high level' actions, while the state-decomposition approach feeds the knowledge of the possible 'low level' agents to a global agent that outputs a full action description.

Thought it was shown that the proposed state-decomposition techniques can be beneficial in terms of sample efficiency, there is an increase in complexity of the design of the agent, compared to simple DQN. The main issue is that the state-transition matrix is not normally known in applications that use reinforcement learning, meaning that if the decomposition is not "intuitive" as in this example, this would have to be estimated before being able to apply state-decomposition. Additionally, there are cases in which there might be different types of possible decompositions and there is no systematic way to predict which one is most efficient. Lastly, while it was shown that the state-decomposition techniques achieve performance gain in certain environment, this is not guaranteed in general, thus the extra design effort might not be rewarded by increased efficiency in some cases. A list of possible extensions of this project, some of which address these issues, is given in Section 7.3.

Chapter 7

Conclusions and further work

7.1 Key findings

Different state-decomposition architectures have been proposed and tested in various environments. It has been found that the initially proposed architecture which is trained in two stages fails to transfer knowledge from the first to the second stage. The most positive results have been obtained with the 'DeltaSwitch' architecture with $\delta = 0$, which simply assigns a different neural network to each of the decomposed subspaces, but allows to train each neural network using target values generated by the other neural networks. This obtains a performance gain over DQN in various environments, as it experiences higher rewards during training, though not necessarily reaching convergence faster. It has been identified that the most likely way in which 'DeltaSwitch' improves performance over DQN is by simplifying the functions that the neural networks aim to learn.

Two transfer learning techniques have also been tested, in order to gain insight on what could be possible techniques to use in two-stage state-decomposition architectures. One of them uses a neural network to combine the networks of the subspaces, while the other combines them using the idea of a progressive neural network. The pretrained networks have been successfully used to provide a good initialisation in the training environment, however, the rate of convergence of these techniques is similar to DQN, meaning that "forgetting" occurs and that knowledge is not successfully transferred from the pretrained networks to the final policy. It was attempted to apply these techniques to a two-stage decomposition approach, however this proves to be difficult as a good initialisation of the second stage requires a long training time in the first stage.

After testing 'DeltaSwitch' in different environments it has also been found that different decomposition types perform better in different scenarios. For example, in 'GridEps2', the best performance is achieved by decomposing by position, which is different from the initially proposed type of decomposition that relies applying a threshold to transition probabilities. This suggests that the decomposition do not necessarily have to separate semi-independent problems to provide a performance gain, as long as it simplifies the functions that the neural networks have to learn.

The two transfer learning techniques inspired two other techniques having only one-stage. These are basically the transfer learning techniques in which the pretrained networks are not pretrained, but are trained simultaneously with the combining network. The architecture using a simple combining

network does not present any performance gain over DQN, while the one using a progressive network architecture shows performance gain in some environments.

It has thus been found that the state-decomposition technique can lead to "cheaper" training, meaning that the rewards are higher during the training process. The 'DeltaSwitch' agent shows the best performance in the tests, while the progressive network architecture shows slightly lower performance but it can be used with a wider range and nature of decompositions. A possible practical application of the proposed method is in the allocation of tasks to interconnected data-centres. Performance gain was only achieved using one stage approaches, however, two-stage approaches could be successful if a way to successfully transfer knowledge to the second stage is determined.

7.2 Discussion

One limitation of this project is that due to timing and computational constraints, the state-decomposition technique has only been tested in relatively simple environments with relatively small state spaces. As this technique aims to aid in scenarios with very large state spaces, it would be useful to test it in such environments. However, the performance gain of state-decomposition could increase further for very large state spaces. The reason is that the environments used in this project are small enough so that all states were likely to be visited multiple times and the neural networks approximating Q-values would not have to rely on generalisation. However, the generalisation properties of neural networks are increasingly important as the size of the state space increases, as it becomes harder to visit all the states during training. As the state-decomposition methods train a sub-agent for each subspace, these agents only have to generalise over the smaller subspaces. It might be that this effect outweighs the fact that there are fewer samples for each subspace than for the whole state space, leading to better generalisation and thus faster convergence and even better performance improvement.

Additionally, for the tested scenarios a higher performance has been achieved using bigger neural networks, due to the small size of the state spaces which leads to better learning with massive neural networks that simply memorize values for each possible state, thus overfitting. However, in scenarios with large state spaces that rely on generalisation, overfitting is very harmful. The state-decomposition technique, by simplifying the functions that the neural networks have to learn, could allow the use of smaller neural networks that are less likely to overfit, leading to better performance in environment with very large state spaces.

A limitation of the state-decomposition method, at least in the explored versions, is that it fails to exploit similarities between subspaces. This is because while dividing the state space into smaller subspaces, it also reduces the number of available training samples for each subspace, as each sample is only used by one subspace. However, some of these samples could provide information for many similar subspaces. For example, in the 'Grid' environment, learning to navigate to one destination clearly has similarities with learning to navigate to the other destination. A more efficient technique would be able to identify these similarities and further improve sample efficiency.

Another problem is that while it has been shown that the state-decomposition method can lead to performance gains in certain environments, using certain network architectures and hyperparam-

eters, this does not necessarily imply that this would also occur with different design choices. For example, all experiments shown in this report used an ϵ -greedy policy and $\gamma = 0.95$, however, different values of these might affect the results.

A limitation of the experimental technique is that some of the results were obtained using a relatively small number of samples (ranging from 20 to 100), which was not always consistent due to timing and computational power constraints. However, the results are considered statistically relevant and it is unlikely that running further trials would have affected to conclusions of this report.

7.3 Further work

Finding the state-transition matrix The first step in the originally proposed state-decomposition method is to obtain the state-transition matrix, as this is usually unknown. A possible extension would be to investigate ways to efficiently estimate this for the purposes of state-decomposition. This could be done by interacting with the environment and sampling the $\langle S, A, R, S' \rangle$ tuples that are generated. Since the state-transition probabilities are dependent on the action, the state-transition matrix is dependent on the chosen policy. As this is conducted before the training, when we have no information about the optimal policy, we could choose to use the most general possible policy to obtain the state-transition table, that is, the actions are all chosen with equal probability independently of the current state. The number of iterations required to construct the state-transition matrix is dependant on the number of states, actions, the probability of visiting each different state and the required level of accuracy of the state-transition matrix estimate, which also would need to be determined.

Compression and embeddings of states The state-decomposition algorithm divides the original state space of the system into smaller groups of states. It can be supposed that the variance of the states within these subspaces is smaller than in the original state space. Thus, the states that are input in the subspaces' agents could be successfully compressed into a more compact representation which could positively affect the performance of the algorithm. More generally, it could be beneficial to learn custom embeddings for the states of each subspace, having the aim of better generalisation rather than state space reduction.

Applying state-decomposition to continuous state spaces The state-decomposition method is only applicable to discrete state spaces. A possible extension would investigate methods to decompose continuous state spaces. A state-transition matrix could be obtained by discretising the states and the state space would be decomposed accordingly, while the agents could still be fed the original continuous states as inputs.

Different types of decomposition This project mostly focuses on decomposition using transition probabilities, however this is not the only possible kind of decomposition, as seen in the data centres example. Example of different decompositions are:

- **Applying threshold to transitions between whole subspaces** The original state-decomposition method consists of temporarily ignoring improbable state transitions by applying a threshold to the state-transition matrix. There are scenarios in which groups of states are connected by state-transitions that have a high probability but it would still be useful to consider them as

separate subspaces. Consider for example Figure-7.1, where an MDP is composed of nearly independent subspaces that are only connected by the blue transitions. In such a scenario it could be useful to determine the subspaces by limiting the number of possible non-zero transition probabilities between any two subspaces rather than applying a probability threshold to every single transition. Alternatively, we could apply a threshold to transition probabilities between whole subspaces instead of single states.

- **Decomposing along the length of the state vector** This could be useful in a case like the example with data-centres, in which different transitions are only likely to affect different sections of the input state's representation vector. Thus, different neural networks could deal with each of these sections of the vector separately.

Sharing information between sub-agents Consider a system composed of many identical sub-systems. The proposed decomposition method would be able to exploit the structure of the system to learn more efficiently. However, in this case, it would be possible to also exploit the similarity of the sub-systems in order to be even more efficient. This is because in this case, the sub-agents would basically be identical, meaning that the samples belonging to one subspace could also be used to train the agents of the other subspaces. This is an extreme scenario, however, there are many environments in which the sub-systems are similar if not identical. Studying ways to share learned information between sub-agents to exploit the similarity of the subspaces could further increase the sample efficiency of training.

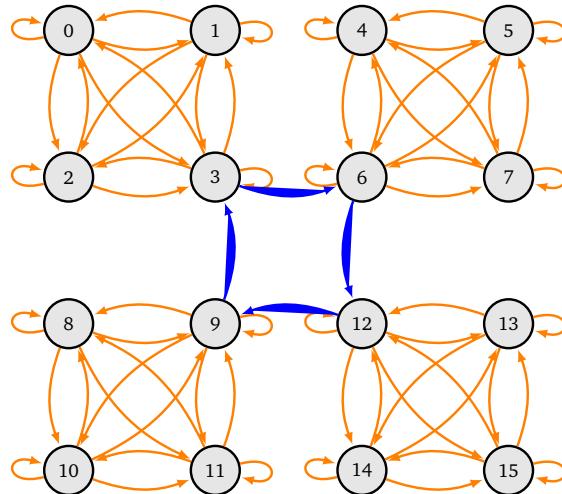


Figure 7.1: MDP composed of smaller subspaces of states connected by few but probable transitions.

Using state-decomposition for automatic sub-goal discovery The idea of state-decomposition, especially when implemented by considering transitions between whole subspaces could be applied to the automatic discovery of sub-goals in hierarchical reinforcement learning. A similar idea has been adopted by [40] and [15], in which the graph of the MDP is divided into subspaces by cutting through the least possible number of state-transitions. The intuition is that sub-goals, such as driving towards the passenger in "Taxi-v3" or driving to the destination after picking it up, are often connected by bottleneck states, such as states 3, 6, 9 and 12 in Figure-7.1.

Proving mathematical bounds It could be useful to derive mathematical bounds for the rate of convergence of different state-decomposition algorithm. This could provide theoretical guidance for more efficient experimentation of new techniques.

Applying state-decomposition to policy gradient algorithms The state-decomposition method was mostly applied in this project to methods based on deep Q-learning. However, state-decomposition is not limited to DQN and it could for example be applied to policy gradient methods.

Bibliography

- [1] Cloud tensor processing units (TPUs). pages 17
- [2] Google AI blog: Scalable deep reinforcement learning for robotic manipulation. pages 2
- [3] Markov decision process. Page Version ID: 1002665350. pages 6
- [4] OpenAI gym: the CartPole-v0 environment. pages 18
- [5] TensorFlow. pages 17
- [6] Mehran Asadi and Manfred Huber. State space reduction for hierarchical reinforcement learning. 2004. pages 11
- [7] Ravi B Kiran, Ibrahim Sobh, Talpaert Victor, Patrick Mannion, Ahmad A. Al Sallab, Senthil Yogamani, and Patrick Pérez. Deep reinforcement learning for autonomous driving: A survey. *arXiv.org*, Feb 2020. pages 2
- [8] Andrew G. Barto and Sridhar Mahadevan. Recent advances in hierarchical reinforcement learning. 13(1):41–77. pages 11
- [9] R. Bellman. Dynamic programming. 153(3731):34–37. pages 7
- [10] Thomas L. Dean, Robert Givan, and Sonia Leach. Model Reduction Techniques for Computing Approximately Optimal Solutions for Markov Decision Processes. *arXiv e-prints*, page arXiv:1302.1533, February 2013. pages 11
- [11] Thomas G. Dietterich. Hierarchical reinforcement learning with the MAXQ value function decomposition. pages 11
- [12] Chris Gaskett, David Wettergreen, and Alexander Zelinsky. Q-learning in continuous state and action spaces. In Norman Foo, editor, *Advanced Topics in Artificial Intelligence*, volume 1747, pages 417–428. Springer Berlin Heidelberg. Series Title: Lecture Notes in Computer Science. pages 10
- [13] David Ha and Jürgen Schmidhuber. World Models. *arXiv e-prints*, page arXiv:1803.10122, March 2018. pages 11
- [14] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. pages 23
- [15] Ishai Menache, Shie Mannor, and Nahum Shimkin. Q-cut—dynamic discovery of sub-goals in reinforcement learning. In Tapio Elomaa, Heikki Mannila, and Hannu Toivonen, editors, *Machine Learning: ECML 2002*, Lecture Notes in Computer Science, pages 295–306. Springer. pages 55
- [16] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin Riedmiller. Playing Atari with Deep Reinforcement Learning. *arXiv e-prints*, page arXiv:1312.5602, December 2013. pages 10, 11
- [17] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dharshan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. Human-level control through deep reinforcement learning. 518(7540):529–533. pages 10

- [18] Thomas M. Moerland, Joost Broekens, and Catholijn M. Jonker. Model-based reinforcement learning: A survey. pages 11
- [19] OpenAI. Gym: A toolkit for developing and comparing reinforcement learning algorithms. pages 18
- [20] Sinno Jialin Pan and Qiang Yang. A survey on transfer learning. 22(10):1345–1359. Conference Name: IEEE Transactions on Knowledge and Data Engineering. pages 39
- [21] Ronald Parr and Stuart Russell. Reinforcement learning with hierarchies of machines. page 7. pages 11
- [22] Romain Paulus, Caiming Xiong, and Richard Socher. A deep reinforced model for abstractive summarization, November 2017. pages 2
- [23] Paul J. Pritz, Liang Ma, and Kin K. Leung. Jointly-Trained State-Action Embedding for Efficient Reinforcement Learning. *arXiv e-prints*, page arXiv:2010.04444, October 2020. pages 11
- [24] Andrei A. Rusu, Neil C. Rabinowitz, Guillaume Desjardins, Hubert Soyer, James Kirkpatrick, Koray Kavukcuoglu, Razvan Pascanu, and Raia Hadsell. Progressive neural networks. pages 43, 44
- [25] SAGAR SHARMA. What the hell is perceptron? pages 8
- [26] David Silver, Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Matthew Lai, Arthur Guez, Marc Lanctot, Laurent Sifre, Dharshan Kumaran, Thore Graepel, Timothy Lillicrap, Karen Simonyan, and Demis Hassabis. Mastering chess and shogi by self-play with a general reinforcement learning algorithm. pages 2, 11
- [27] Richard S. Sutton. Dyna, an integrated architecture for learning, planning, and reacting. 2(4):160–163. pages 11
- [28] Richard S. Sutton and Andrew G. Barto. *Reinforcement learning: an introduction*. The MIT Press, 2018. pages 2, 4
- [29] Richard S. Sutton, Doina Precup, and Satinder Singh. Between MDPs and semi-MDPs: A framework for temporal abstraction in reinforcement learning. 112(1):181–211. pages 11
- [30] Matthew E Taylor and Peter Stone. Transfer learning for reinforcement learning domains: A survey. page 53. pages 39
- [31] Keras Team. Simple. flexible. powerful. Accessed: 2021-01-23. pages 17
- [32] Hado van Hasselt, Arthur Guez, and David Silver. Deep reinforcement learning with double q-learning. pages 10
- [33] Ziyu Wang, Tom Schaul, Matteo Hessel, Hado van Hasselt, Marc Lanctot, and Nando de Freitas. Dueling network architectures for deep reinforcement learning. page 9. pages 10
- [34] Christopher J. C. H. Watkins and Peter Dayan. Q-learning. 8(3):279–292. pages 10
- [35] Christopher John Cornish Hellaby. Watkins. *Learning from delayed rewards*. PhD thesis, University of Cambridge, 1989. pages 9
- [36] Centrum Wiskunde. *DoubleQ-learning Hado van Hasselt Multi-agent and Adaptive Computation Group*. pages 10
- [37] Chao Yu, Jiming Liu, and Shamim Nemati. Reinforcement learning in healthcare: A survey. pages 2
- [38] Ziyao Zhang, Liang Ma, Kin K. Leung, Leandros Tassiulas, and Jeremy Tucker. Q-placement: Reinforcement-learning-based service placement in software-defined networks. *2018 IEEE 38th International Conference on Distributed Computing Systems (ICDCS)*, 2018. pages 2
- [39] Ziyao Zhang, Liang Ma, Konstantinos Poularakis, Kin K. Leung, and Lingfei Wu. Dq scheduler: Deep reinforcement learning based controller synchronization in distributed sdn. *ICC 2019 - 2019 IEEE International Conference on Communications (ICC)*, 2019. pages 2

- [40] Özgür Şimşek, Alicia P. Wolfe, and Andrew G. Barto. Identifying useful subgoals in reinforcement learning by local graph partitioning. In *Proceedings of the 22nd international conference on Machine learning*, ICML '05, pages 816–823. Association for Computing Machinery. pages 55

Appendix A

Performance graphs

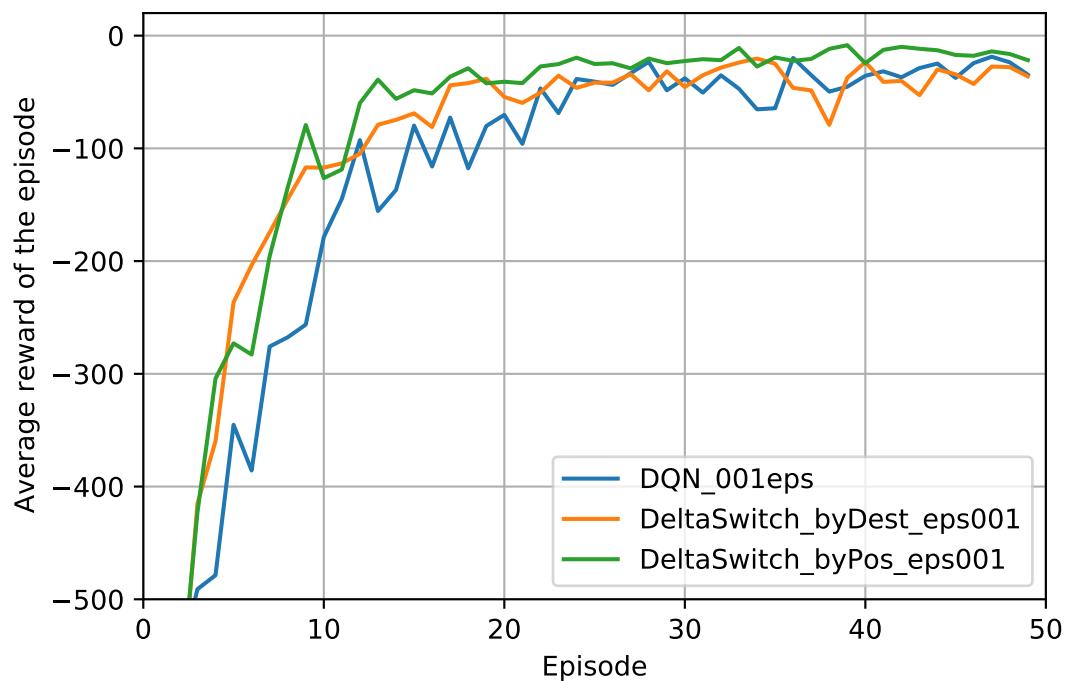


Figure A.1: Performance of DQN and 'DeltaSwitch' in 'GridEps2'.

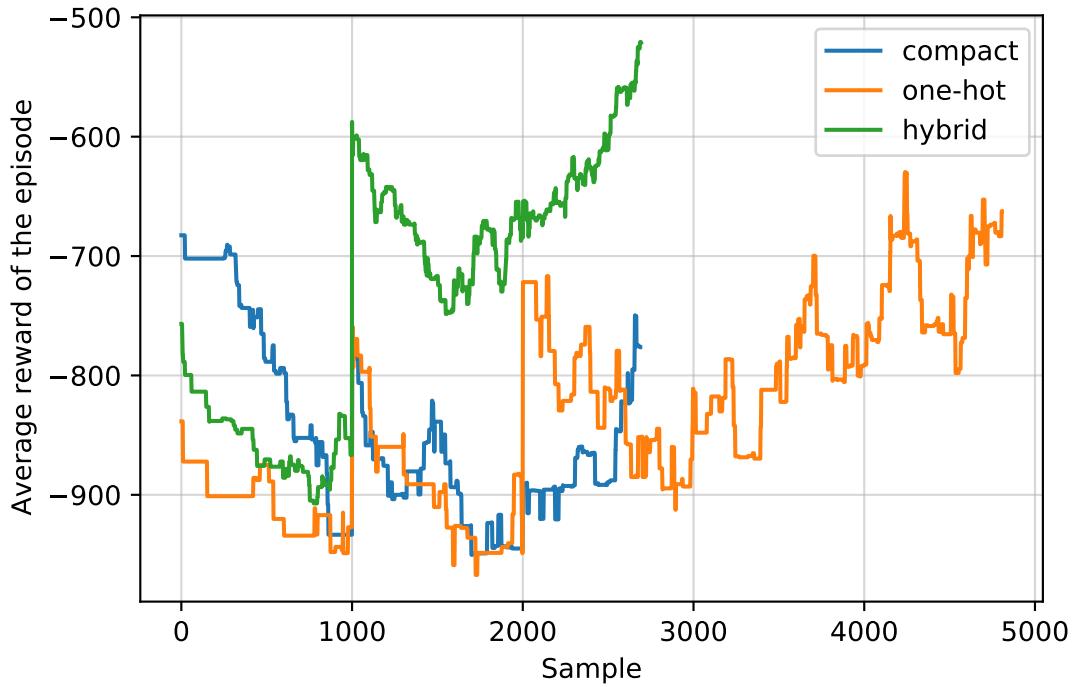


Figure A.2: Effect of different state encoding methods on the performance of the baseline DQ-Nagent in the 'Grid' environment.

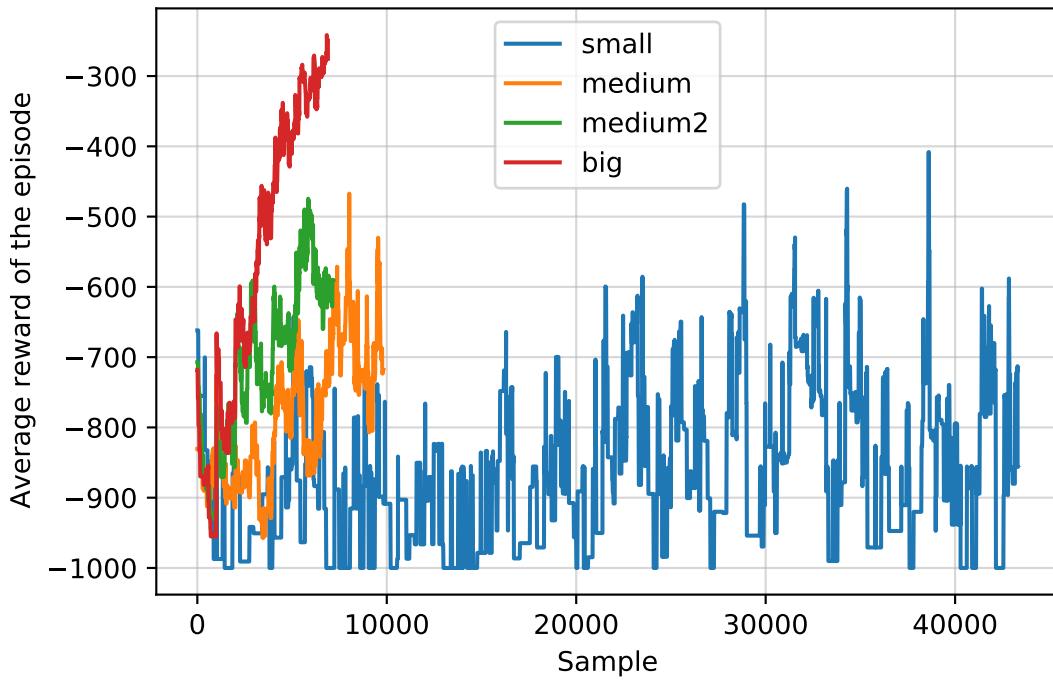


Figure A.3: Performance of DQN with NNs of different sizes in the Taxi2 environment. 'medium2' has twice the number of trainable weights of 'medium'.

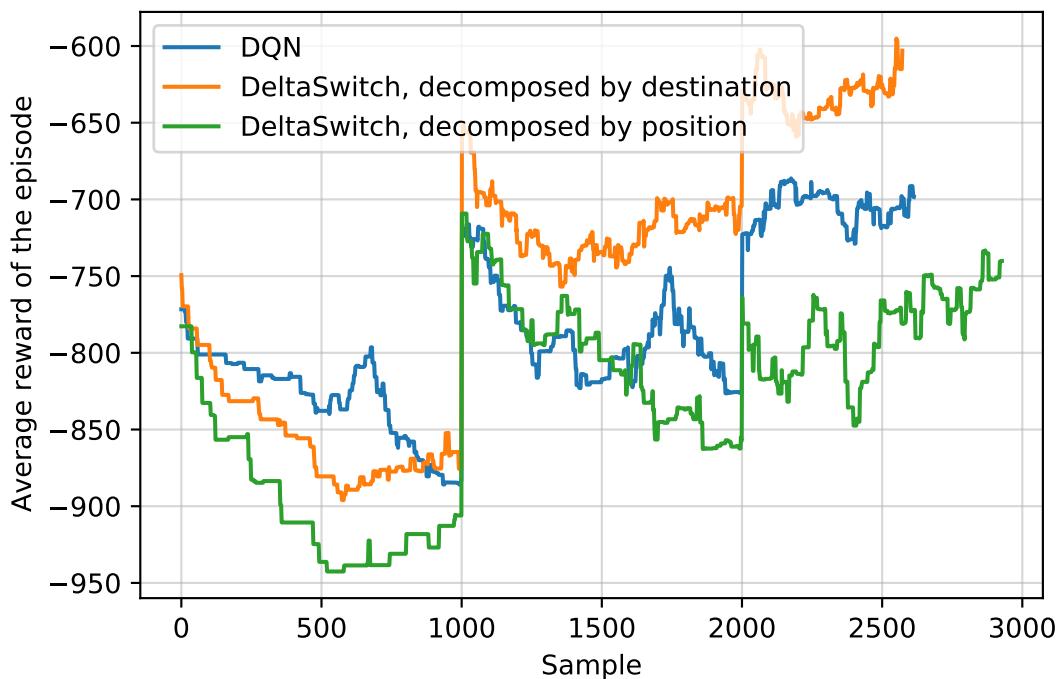


Figure A.4: Comparison of DQN and DeltaSwitch with $\delta = 0$ in the Grid environment.

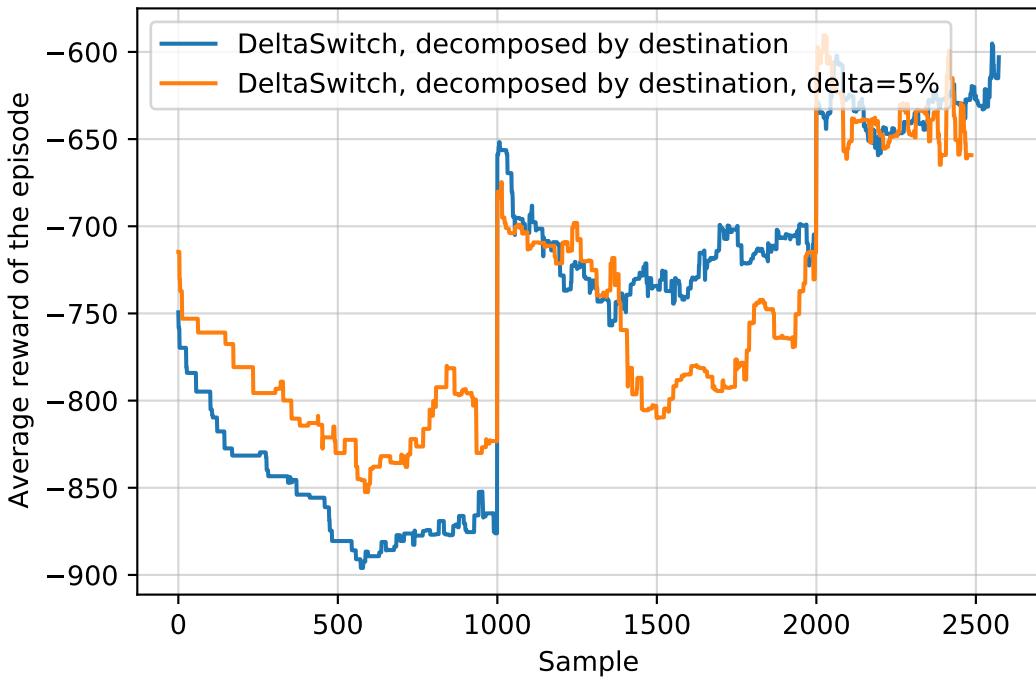
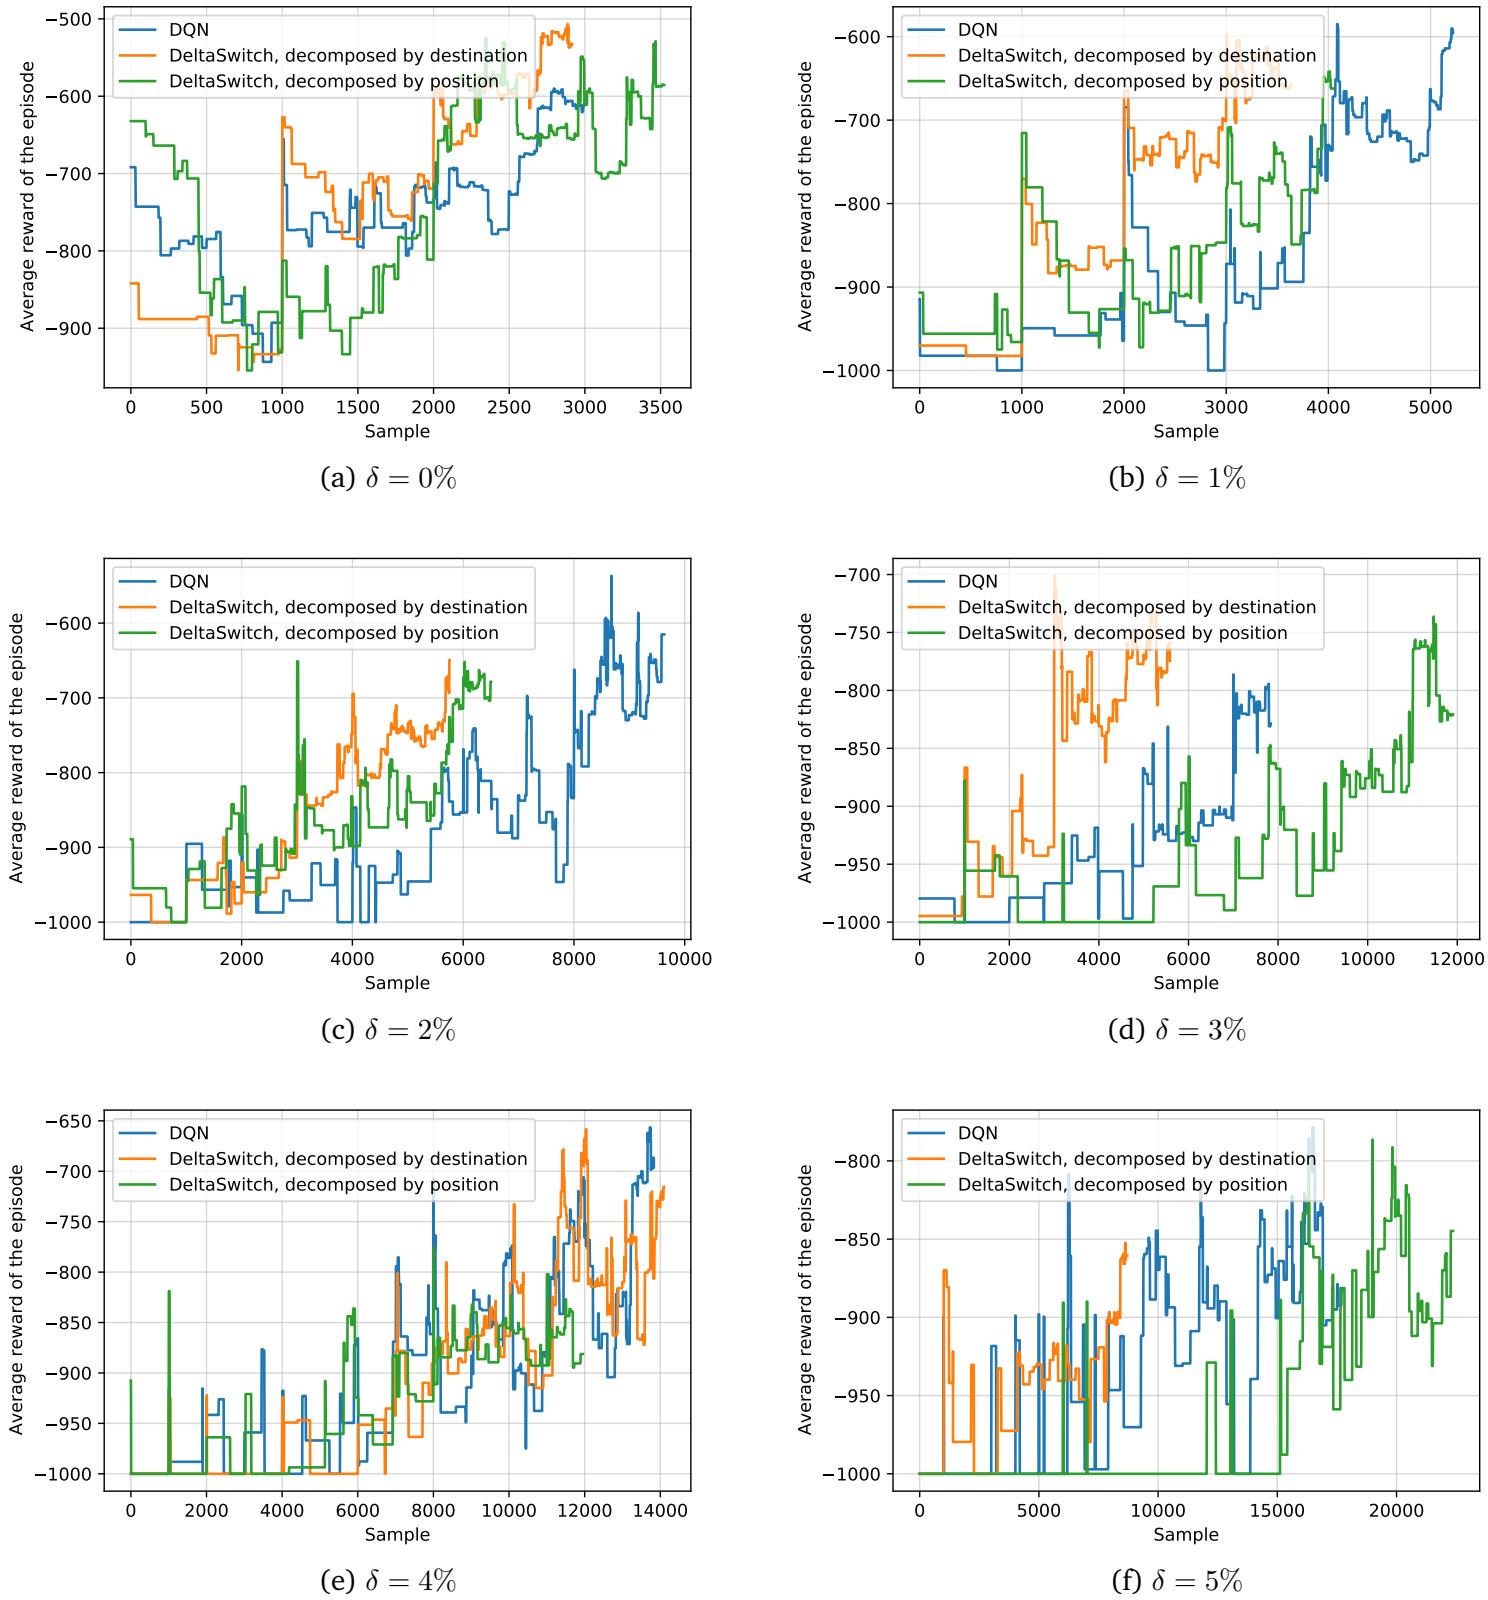


Figure A.5: Comparison of DeltaSwitch with different δ .



Figure A.7: Effect of reducing the input size.


 Figure A.6: DeltaSwitch vs DQN in 'GridEps' with $0 \leq \epsilon \leq 5\%$.

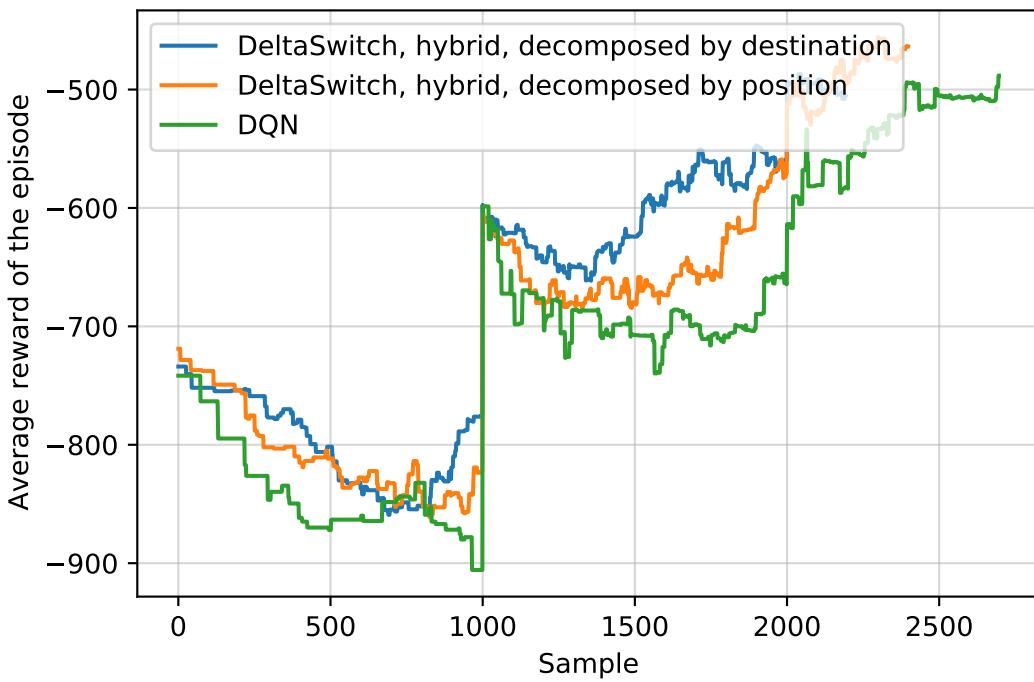


Figure A.8: Performance of DQN and DeltaSwitch with different decomposition in the 'Grid' environment with wall.

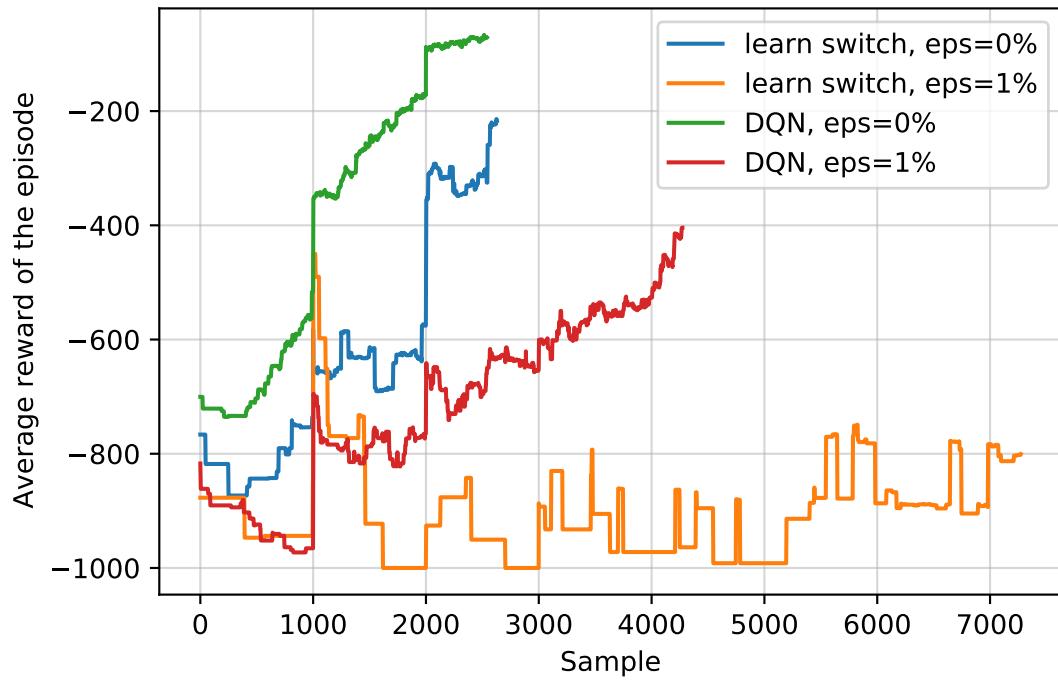


Figure A.9: Performance of combining 'NN0' and 'NN1' using another NN that is also fed the current state as input in a 'GridEps' environment with $\epsilon = 0\%, 1\%$.

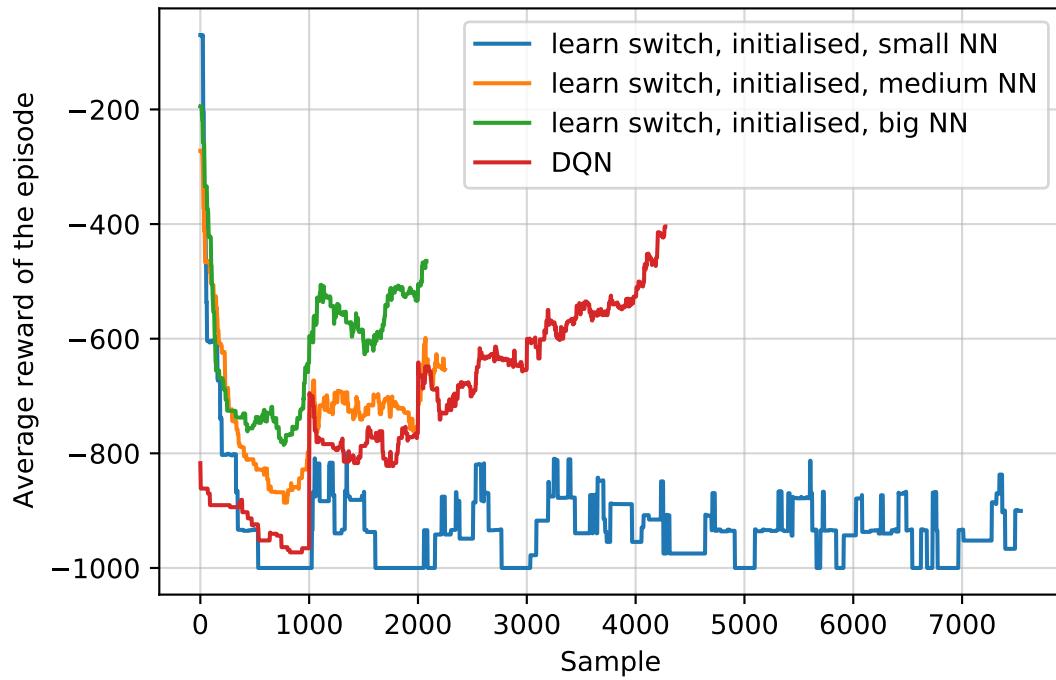


Figure A.10: Performance of combining 'NN0' and 'NN1' with another NN when initialising the combined network to behave as the original 'DeltaSwitch'. Tested with different network sizes

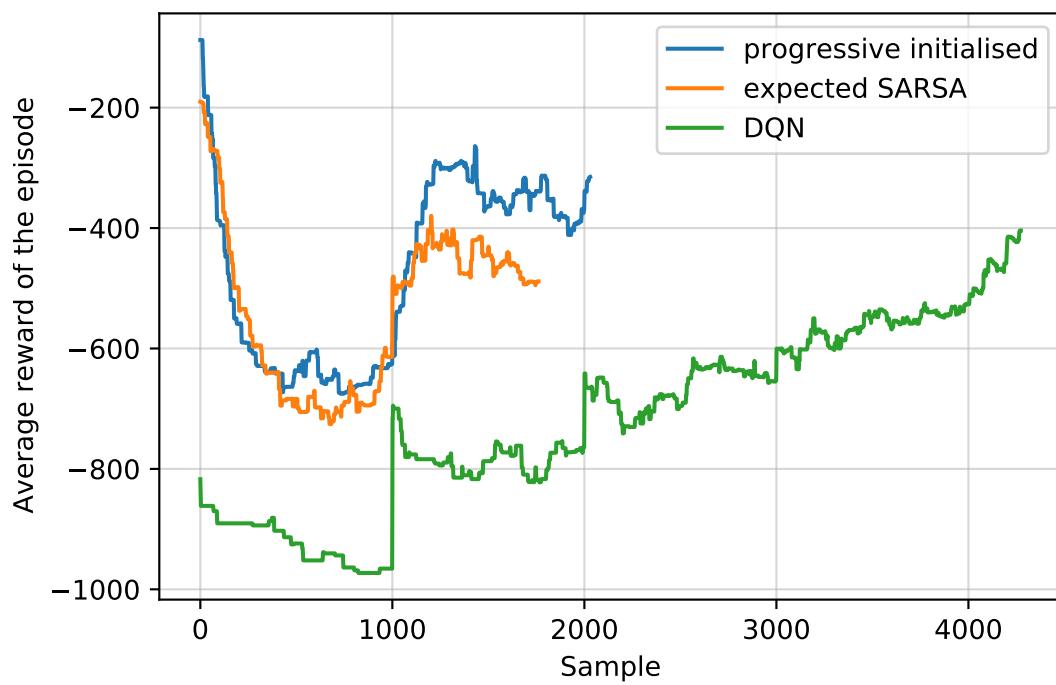


Figure A.11: Comparison of using Expected SARSA after combining, progressive combined network and standard DQN.

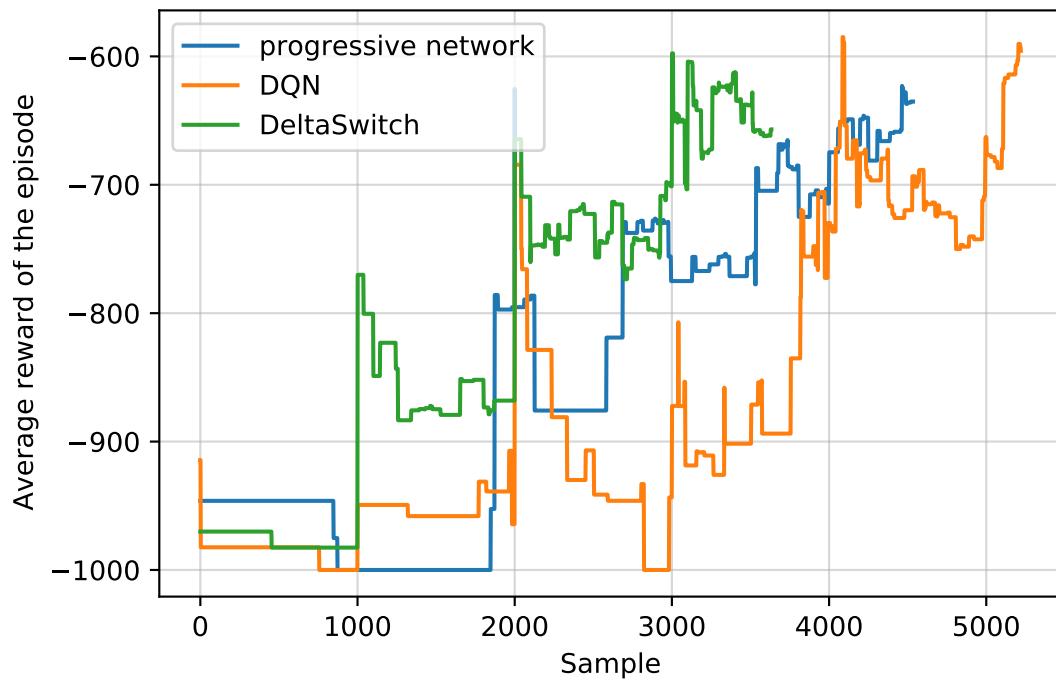


Figure A.12: Performance of progressive network and combining network when 'NN0' and 'NN1' are not pre-trained. Standard DQN and DeltaSwitch are shown for comparison.